

```

import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns

import importlib

from importlib import reload

plt=reload(plt)

data=pd.read_csv('logistic_regression.csv')

data.head()

```

	loan_amnt	term	int_rate	installment	grade	sub_grade	\
0	10000	36 months	11.44	329.48	B	B4	
1	8000	36 months	11.99	265.68	B	B5	
2	15600	36 months	10.49	506.97	B	B3	
3	7200	36 months	6.49	220.65	A	A2	
4	24375	60 months	17.27	609.33	C	C5	

	emp_title	emp_length	home_ownership	annual_inc	...
0	Marketing	10+ years	RENT	117000.0	...
1	Credit analyst	4 years	MORTGAGE	65000.0	...
2	Statistician	< 1 year	RENT	43057.0	...
3	Client Advocate	6 years	RENT	54000.0	...
4	Destiny Management Inc.	9 years	MORTGAGE	55000.0	...

	open_acc	pub_rec	revol_bal	revol_util	total_acc	initial_list_status
0	16	0	36369	41.8	25	w
1	17	0	20131	53.3	27	f
2	13	0	11987	92.2	26	f
3	6	0	5472	21.5	13	f
4	13	0	24584	69.8	43	f

	application_type	mort_acc	pub_rec_bankruptcies	\
0	INDIVIDUAL	0.0	0.0	
1	INDIVIDUAL	3.0	0.0	
2	INDIVIDUAL	0.0	0.0	

```
3      INDIVIDUAL      0.0      0.0
4      INDIVIDUAL      1.0      0.0
```

```
                                address
0      0174 Michelle Gateway\nMendozaberg, OK 22690
1      1076 Carney Fort Apt. 347\nLoganmouth, SD 05113
2      87025 Mark Dale Apt. 269\nNew Sabrina, WV 05113
3      823 Reid Ford\nDelacruzside, MA 00813
4      679 Luna Roads\nGreggshire, VA 11650
```

```
[5 rows x 27 columns]
```

```
data.shape
```

```
(396030, 27)
```

```
data.columns
```

```
Index(['loan_amnt', 'term', 'int_rate', 'installment', 'grade',
      'sub_grade',
      'emp_title', 'emp_length', 'home_ownership', 'annual_inc',
      'verification_status', 'issue_d', 'loan_status', 'purpose',
      'title',
      'dti', 'earliest_cr_line', 'open_acc', 'pub_rec', 'revol_bal',
      'revol_util', 'total_acc', 'initial_list_status',
      'application_type',
      'mort_acc', 'pub_rec_bankruptcies', 'address'],
      dtype='object')
```

```
data.dtypes
```

```
loan_amnt      int64
term            object
int_rate       float64
installment    float64
grade          object
sub_grade      object
emp_title      object
emp_length     object
home_ownership object
annual_inc     float64
verification_status object
issue_d        object
loan_status    object
purpose        object
title          object
dti            float64
earliest_cr_line object
open_acc       int64
pub_rec        int64
revol_bal      int64
```


```
revol_util          float64
total_acc           int64
initial_list_status object
application_type    object
mort_acc            float64
pub_rec_bankruptcies float64
address             object
dtype: object
```

```
data.isnull().sum()
```

```
loan_amnt          0
term               0
int_rate           0
installment        0
grade              0
sub_grade          0
emp_title          22927
emp_length         18301
home_ownership     0
annual_inc         0
verification_status 0
issue_d            0
loan_status        0
purpose            0
title              1756
dti                0
earliest_cr_line   0
open_acc           0
pub_rec            0
revol_bal          0
revol_util         276
total_acc          0
initial_list_status 0
application_type    0
mort_acc           37795
pub_rec_bankruptcies 535
address            0
dtype: int64
```


```
data.loan_amnt.hist()
```

```
<Axes: >
```



```
data.annual_inc.hist() #This shows the presence of a outliers
```

```
<Axes: >
```



```
data.annual_inc.describe()
count      3.960300e+05
mean       7.420318e+04
std        6.163762e+04
min        0.000000e+00
25%        4.500000e+04
50%        6.400000e+04
75%        9.000000e+04
max        8.706582e+06
Name: annual_inc, dtype: float64
```

```
data.term.value_counts()
term
36 months    302005
60 months     94025
Name: count, dtype: int64
```

```
data.term=data.term.str.split(expand=True)[0]
```

```
data.term.value_counts()
```

```

term
36      302005
60       94025
Name: count, dtype: int64


data.int_rate.describe()

count      396030.000000
mean        13.639400
std         4.472157
min         5.320000
25%        10.490000
50%        13.330000
75%        16.490000
max        30.990000
Name: int_rate, dtype: float64

data.int_rate.hist()

<Axes: >

```




```

data.grade.value_counts().sort_index().plot(kind='bar')

<Axes: xlabel='grade'>

```



```
data.grade.value_counts()
```

```
grade
```

```
B    116018
```

```
C    105987
```

```
A     64187
```

```
D     63524
```

```
E     31488
```

```
F     11772
```

```
G       3054
```

```
Name: count, dtype: int64
```

```
sorted(data.sub_grade.unique()) #assuming A1 is better than F5 because  
this is a grade and ordinal
```

```
['A1',  
'A2',  
'A3',  
'A4',  
'A5',  
'B1',  
'B2',  
'B3',
```

```
'B4',  
'B5',  
'C1',  
'C2',  
'C3',  
'C4',  
'C5',  
'D1',  
'D2',  
'D3',  
'D4',  
'D5',  
'E1',  
'E2',  
'E3',  
'E4',  
'E5',  
'F1',  
'F2',  
'F3',  
'F4',  
'F5',  
'G1',  
'G2',  
'G3',  
'G4',  
'G5']
```

```
data.emp_title.value_counts()
```

```
emp_title  
Teacher          4389  
Manager          4250  
Registered Nurse  1856  
RN               1846  
Supervisor       1830  
...  
sikorsky         1  
Postman          1  
McCarthy & Holthus, LLC  1  
jp flooring      1  
Gracon Services, Inc  1  
Name: count, Length: 173103, dtype: int64
```

```
data[data.emp_title.isnull()]
```

```
      loan_amnt term  int_rate  installment grade sub_grade  
emp_title \
```


mort_acc \				
35	88.7	20	f	INDIVIDUAL
5.0				
36	54.6	7	f	INDIVIDUAL
0.0				
40	53.8	27	f	INDIVIDUAL
10.0				
49	93.2	18	w	INDIVIDUAL
6.0				
58	82.4	7	w	INDIVIDUAL
0.0				
...
...				
395946	18.8	21	w	INDIVIDUAL
5.0				
395963	9.7	44	w	INDIVIDUAL
0.0				
395988	32.4	33	w	INDIVIDUAL
1.0				
395999	81.0	24	w	INDIVIDUAL
4.0				
396015	5.8	27	w	INDIVIDUAL
5.0				

pub_rec_bankruptcies \	
35	1.0
36	0.0
40	0.0
49	0.0
58	0.0
...	...
395946	0.0
395963	1.0
395988	0.0
395999	0.0
396015	0.0

address	
35	23617 Michael Viaduct\nWest John, MS 05113
36	361 Erica Forest\nLake Mariaton, TN 30723
40	8268 Reed Gardens Suite 996\nEast Johnmouth, N...
49	84009 David Stream\nSouth Nicolehaven, IL 05113
58	965 Spencer Courts\nPacetown, AZ 00813
...	...
395946	2645 Wayne Street\nMarymouth, HI 22690
395963	8339 Daniel Forges Suite 273\nPort Oscarmouth,...
395988	114 Sonya Pass\nCarlamouth, SD 00813
395999	1314 Bridget Terrace\nRebeccashire, NE 30723
396015	Unit 4067 Box 2110\nDPO AA 05113

[22927 rows x 27 columns]

data[data.emp_length.isnull()]

emp_title \	loan_amnt	term	int_rate	installment	grade	sub_grade
35	5375	36	13.11	181.39	B	B4
NaN						
36	3250	36	16.78	115.52	C	C5
NaN						
49	15000	36	7.89	469.29	A	A5
NaN						
58	10000	36	17.56	359.33	D	D1
NaN						
91	30225	60	18.24	771.47	D	D5
NaN						
...
.						
395946	35000	60	16.20	854.86	C	C4
NaN						
395963	7000	36	20.20	260.86	E	E3
NaN						
395988	35000	60	15.59	843.53	D	D1
NaN						
395999	11125	36	24.11	437.11	F	F2
NaN						
396015	4000	36	9.16	127.50	B	B2
NaN						

emp_title \	emp_length	home_ownership	annual_inc	...	open_acc	pub_rec
35	NaN	RENT	34000.00	...	9	1
14998						
36	NaN	RENT	22500.00	...	7	0
7587						
49	NaN	MORTGAGE	90000.00	...	7	0
8205						
58	NaN	MORTGAGE	32000.00	...	6	0
11615						
91	NaN	MORTGAGE	65800.00	...	11	0
14390						
...
...						
395946	NaN	MORTGAGE	84000.00	...	7	0
4241						
395963	NaN	OWN	32964.00	...	24	1
3236						
395988	NaN	OWN	102396.00	...	15	0
31665						

395999	NaN	MORTGAGE	31789.88	...	8	0
22385						
396015	NaN	MORTGAGE	57400.00	...	12	0
3134						

	revol_util	total_acc	initial_list_status	application_type
mort_acc \				
35	88.7	20	f	INDIVIDUAL
5.0				
36	54.6	7	f	INDIVIDUAL
0.0				
49	93.2	18	w	INDIVIDUAL
6.0				
58	82.4	7	w	INDIVIDUAL
0.0				
91	69.5	31	w	INDIVIDUAL
1.0				
...
...				
395946	18.8	21	w	INDIVIDUAL
5.0				
395963	9.7	44	w	INDIVIDUAL
0.0				
395988	32.4	33	w	INDIVIDUAL
1.0				
395999	81.0	24	w	INDIVIDUAL
4.0				
396015	5.8	27	w	INDIVIDUAL
5.0				

	pub_rec_bankruptcies
35	1.0
36	0.0
49	0.0
58	0.0
91	0.0
...	...
395946	0.0
395963	1.0
395988	0.0
395999	0.0
396015	0.0

	address
35	23617 Michael Viaduct\nWest John, MS 05113
36	361 Erica Forest\nLake Mariaton, TN 30723
49	84009 David Stream\nSouth Nicolehaven, IL 05113
58	965 Spencer Courts\nPacetown, AZ 00813
91	493 Michael Route\nHillfurt, AZ 70466
...	...

```
395946          2645 Wayne Street\nMarymouth, HI 22690
395963  8339 Daniel Forges Suite 273\nPort Oscarmouth,...
395988          114 Sonya Pass\nCarlamouth, SD 00813
395999          1314 Bridget Terrace\nRebeccashire, NE 30723
396015          Unit 4067 Box 2110\nDPO AA 05113
```

```
[18301 rows x 27 columns]
```

```
data.home_ownership.value_counts()
```

```
home_ownership
MORTGAGE      198348
RENT          159790
OWN           37746
OTHER          112
NONE           31
ANY            3
Name: count, dtype: int64
```

```
data.groupby(['home_ownership', 'loan_status'])['loan_status'].size()
```

```
home_ownership  loan_status
ANY             Fully Paid      3
MORTGAGE        Charged Off    33632
                Fully Paid    164716
NONE            Charged Off      7
                Fully Paid     24
OTHER           Charged Off     16
                Fully Paid     96
OWN             Charged Off    7806
                Fully Paid    29940
RENT            Charged Off    36212
                Fully Paid   123578
Name: loan_status, dtype: int64
```

```
ho_fp=data[data['loan_status']=='Fully
Paid'].groupby('home_ownership')['loan_status'].count()
```

```
ho_fp
```

```
home_ownership
ANY            3
MORTGAGE      164716
NONE          24
OTHER         96
OWN          29940
```

```

RENT      123578
Name: loan_status, dtype: int64

ho_tls=data.groupby('home_ownership')['loan_status'].count()

ho_tls


home_ownership
ANY      3
MORTGAGE 198348
NONE     31
OTHER    112
OWN      37746
RENT     159790
Name: loan_status, dtype: int64

perc_default_by_ho=(1-(ho_fp/ho_tls))*100

perc_default_by_ho.sort_values(ascending=False).plot(kind='bar')
plt.title('People who are staying on rent or do not own a house have
maximum chances of default')
plt.ylabel('Default rate')
plt.show()


```

People who are staying on rent or do not own a house have maximum chances of default




```
sns.scatterplot(x='annual_inc', y='loan_amnt', hue='loan_status',  
data=data) #The loans are more for people who do not have great  
incomes  
#So there is an opportunity to target more high value customers for  
loans
```

```
<Axes: xlabel='annual_inc', ylabel='loan_amnt'>
```



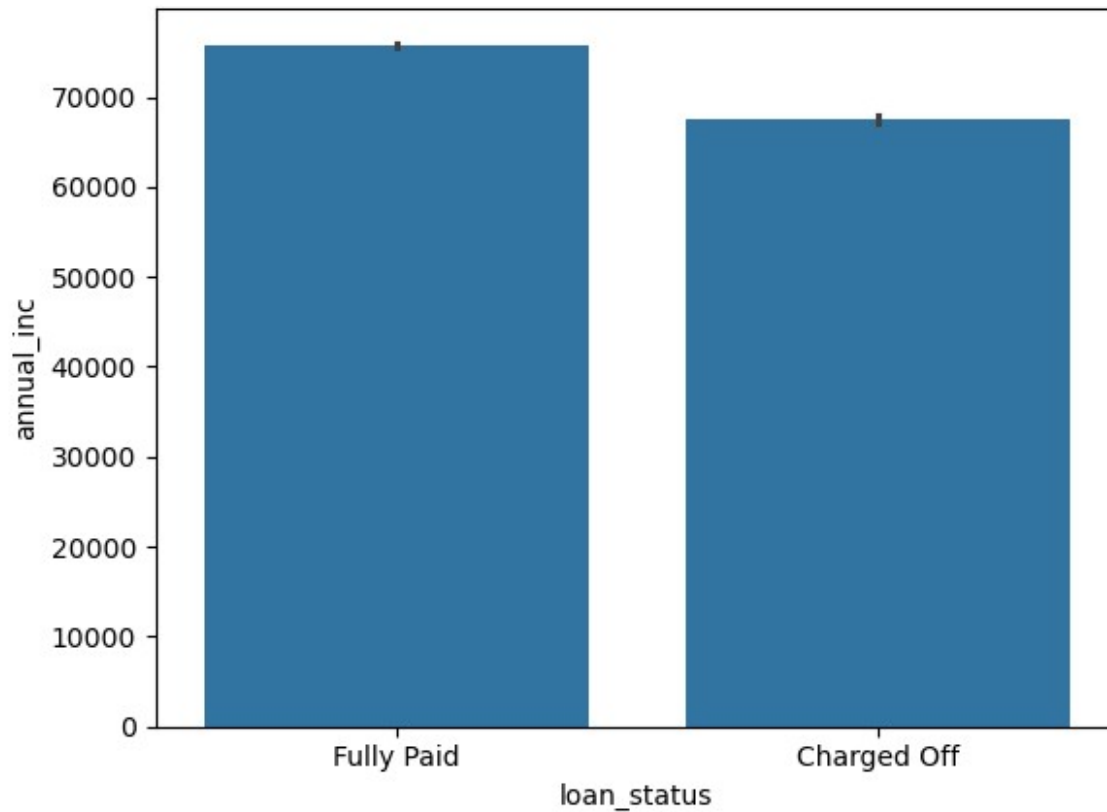
```
sns.barplot(x='loan_status', y='loan_amnt', data=data) #people who  
have charged off took higher loans on average
```

```
<Axes: xlabel='loan_status', ylabel='loan_amnt'>
```



```
sns.barplot(y='annual_inc',x='loan_status', data=data) #People who charged off had lower incomes on average
```

```
<Axes: xlabel='loan_status', ylabel='annual_inc'>
```

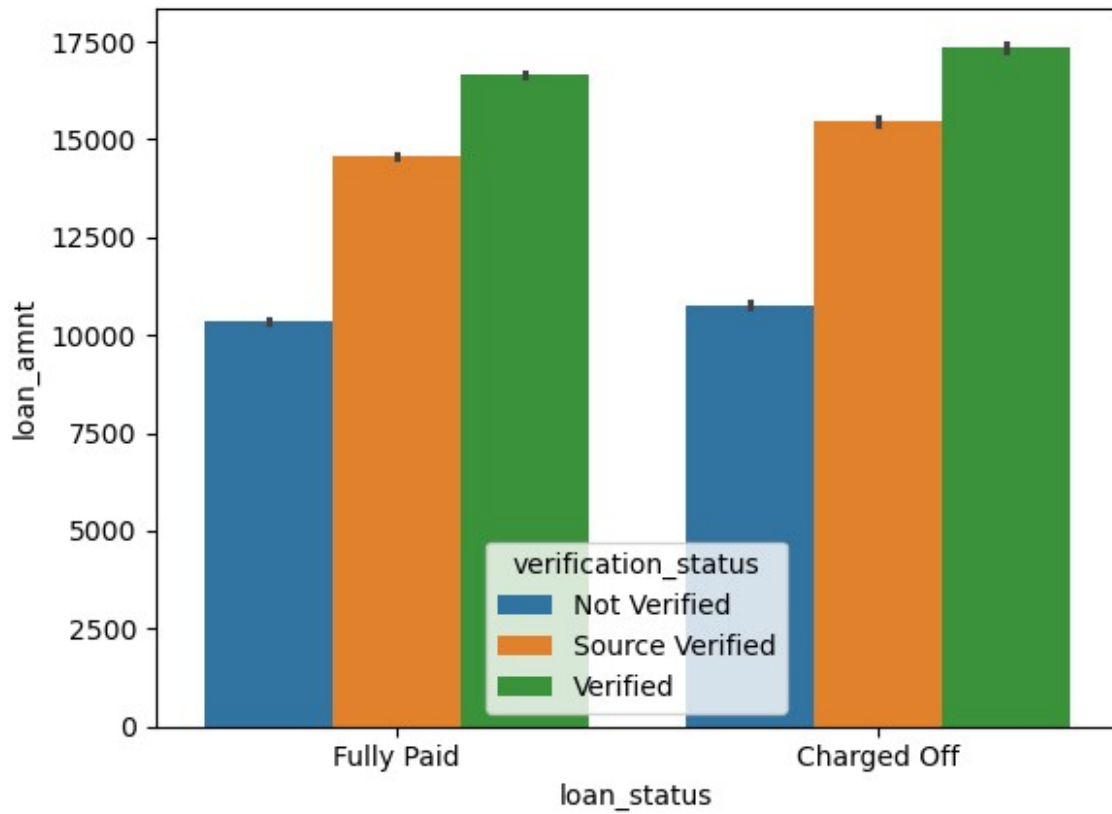



```
data.verification_status.value_counts()
```

```
verification_status
Verified          139563
Source Verified   131385
Not Verified      125082
Name: count, dtype: int64
```


```
sns.barplot(x='loan_status', y='loan_amnt', hue='verification_status',
data=data) #verified people get higher loans
#People who took ran away got higher loans than those who paid up on time
```

```
<Axes: xlabel='loan_status', ylabel='loan_amnt'>
```




```
sns.barplot(x='loan_status', y='annual_inc',  
hue='verification_status', data=data) #People who paid up had higher  
incomes than people who did not
```

```
<Axes: xlabel='loan_status', ylabel='annual_inc'>
```



```
sns.countplot(x='grade', hue='loan_status',  
order=['A','B','C','D','E','F','G'],data=data) #Since this depends on  
the count of the data, I am going to take a %
```


```
<Axes: xlabel='grade', ylabel='count'>
```



```
perc_paid_by_grade=100*data[data['loan_status']=='Fully  
Paid'].groupby('grade')['loan_status'].count()/data.groupby('grade')  
['loan_status'].count()
```

```
perc_paid_by_grade.plot(kind='bar') #clearly, having an "A" grade  
ensures that more people who take the loans pay up
```


```
<Axes: xlabel='grade'>
```



```
data.groupby('grade')['loan_amnt'].sum().plot(kind='bar')
plt.title('The share of total loan amount is more for B & C categories
compared to A which is more safer')
plt.show()
```


#More loans are going to b and c grade people when compared to a grade people and this increases the risk
#Therefore, there is an opportunity for the company to reroute most of its loan amount to a grade customers.

The share of total loan amount is more for B & C categories compared to A which is more safer




```
data.grade.value_counts().sort_index().plot(kind='bar')
```

```
<Axes: xlabel='grade'>
```




```
data.groupby('grade')['loan_amnt'].mean().plot(kind='bar')  
#This shows that on an average a higher loan is being provided to  
those who have the highest risk of default.  
#This means the company should lower the risk of losing money, instead  
providing the highest loans to grade a customers
```

```
<Axes: xlabel='grade'>
```



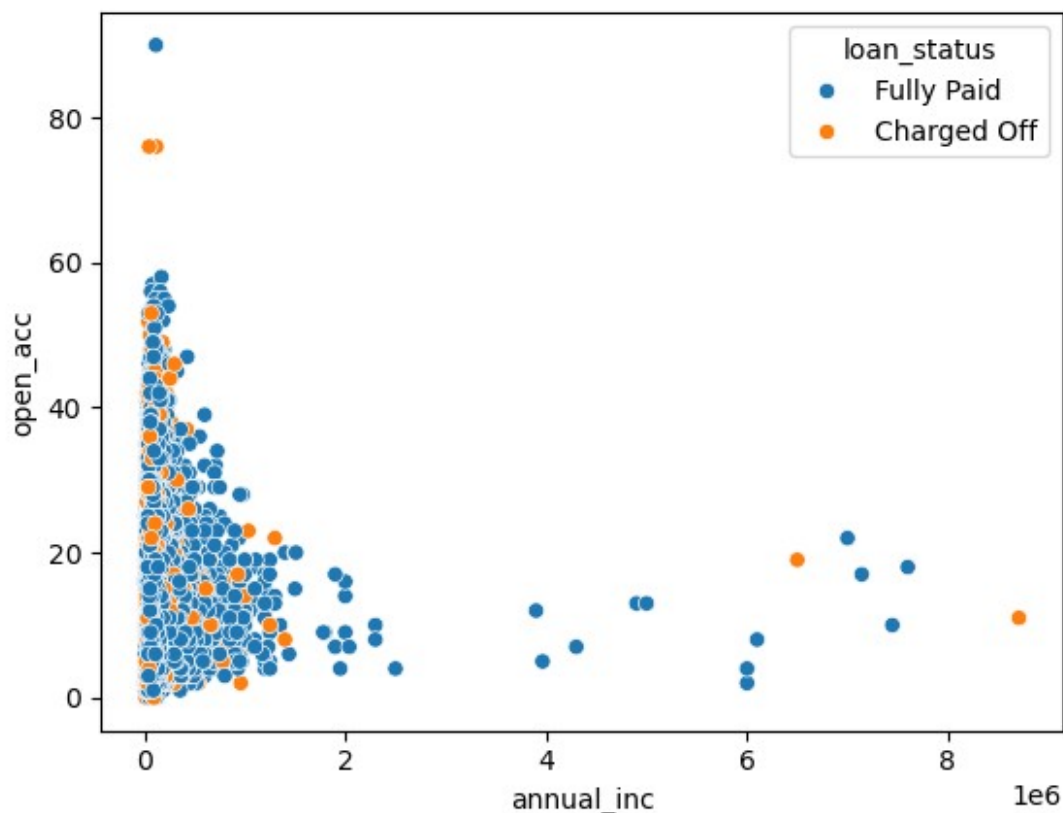
```
sns.scatterplot(x='open_acc', y='loan_amnt', hue='loan_status',  
data=data)
```

```
<Axes: xlabel='open_acc', ylabel='loan_amnt'>
```



```
sns.scatterplot(x='annual_inc', y='open_acc', hue='loan_status',  
data=data)
```

```
<Axes: xlabel='annual_inc', ylabel='open_acc'>
```

```
data.initial_list_status.value_counts()
```

```
initial_list_status
```

```
f    238066
```

```
w    157964
```

```
Name: count, dtype: int64
```

```
data.loan_status.value_counts()
```

```
loan_status
```

```
Fully Paid    318357
```

```
Charged Off    77673
```

```
Name: count, dtype: int64
```

```
data.groupby(['initial_list_status', 'loan_status'])
```

```
['loan_status'].count()
```

```
initial_list_status  loan_status
```

```
f                    Charged Off    45961
```

```
                    Fully Paid    192105
```

```
w                    Charged Off    31712
```

```
                    Fully Paid    126252
```

```
Name: loan_status, dtype: int64
```

```

init_status_default_risk=(data[data['loan_status']=='Charged
Off'].groupby('initial_list_status')['initial_list_status'].count())/
data.groupby('initial_list_status')['initial_list_status'].count()

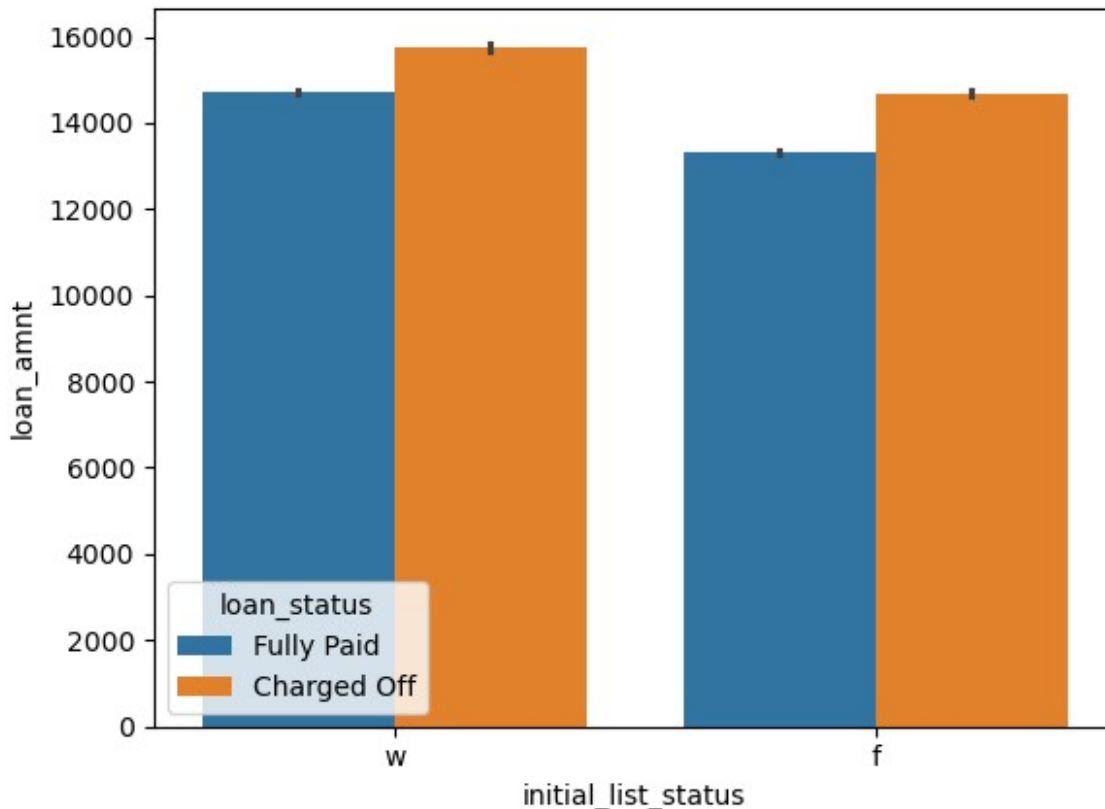
init_status_default_risk #There is not much of a difference here

initial_list_status
f    0.193060
w    0.200755
Name: initial_list_status, dtype: float64

sns.barplot(x='initial_list_status', y='loan_amnt', hue='loan_status',
data=data)

<Axes: xlabel='initial_list_status', ylabel='loan_amnt'>

```




```

sns.barplot(x='initial_list_status', y='annual_inc',
hue='loan_status', data=data)


<Axes: xlabel='initial_list_status', ylabel='annual_inc'>

```




```
sns.scatterplot(x='pub_rec', y='loan_amnt', hue='loan_status',  
data=data)
```

```
<Axes: xlabel='pub_rec', ylabel='loan_amnt'>
```



```
sns.scatterplot(x='pub_rec', y='annual_inc', hue='loan_status',  
data=data)
```

```
<Axes: xlabel='pub_rec', ylabel='annual_inc'>
```



```
data.purpose.value_counts()
```

#this shows that the top 2 categories are debt_consolidation and credit cards, both of which indicate that one loan is funding the recovery of other loans

#This means the company is at high risk from these categories, and should move to give loans in categories where there is security available


purpose	
debt_consolidation	234507
credit_card	83019
home_improvement	24030
other	21185
major_purchase	8790
small_business	5701
car	4697
medical	4196
moving	2854
vacation	2452
house	2201
wedding	1812
renewable_energy	329

```
educational          257
Name: count, dtype: int64
```

```
data.groupby('purpose')
['loan_amnt'].sum().sort_values(ascending=False) #this confirms the
above finding, the loan should be in areas where there is security,
not to repay other risky loans
```

```
purpose
debt_consolidation    3489116875
credit_card           1202306225
home_improvement      339160650
other                 203980925
major_purchase        96065225
small_business        87722650
car                   38615950
medical               37574525
house                 33884825
moving                22473500
wedding               18540025
vacation              15252050
renewable_energy      3076750
educational           1752925
Name: loan_amnt, dtype: int64
```


```
(data[data['loan_status']=='Charged Off'].groupby('purpose')
['purpose'].count()*100/data.groupby('purpose')
['purpose'].count()).sort_values().plot(kind='bar')
plt.title('% risk of default from different purposes')
plt.show()
```



#From the graph above, we see that wedding and educational loans are low risk, yet the company gives fewer loans to them

#Now, is the interest rate correlated with risk?

```
data.groupby('purpose')  
['int_rate'].mean().sort_values().plot(kind='bar')  
plt.title("The interest rates are correlated to risk by category")  
plt.show()
```



#This shows that the bank gives high interest rates to risky categories, and penalizes innocent borrowers to account for those who are defaulting

#The most likely reason is because they do not maintain data at an individual level and adjust it regularly based on each customer's behaviour

#Therefore in this case, I recommend that the bank should implement a blockchain solution so that they can identify who is likely to default rather than just which category is at high risk using aggregate data

```
data.application_type.value_counts()
```

```
application_type
INDIVIDUAL      395319
JOINT            425
```




```
DIRECT_PAY      286
Name: count, dtype: int64
```

```
data.groupby('application_type')['loan_status'].value_counts() #This shows that we should look only at individual borrowers only
```

```
application_type  loan_status
DIRECT_PAY        Fully Paid      184
                  Charged Off     102
INDIVIDUAL        Fully Paid    317802
                  Charged Off    77517
JOINT             Fully Paid      371
                  Charged Off      54
Name: count, dtype: int64
```

```
sns.countplot(x='application_type', hue='loan_status', data=data)
<Axes: xlabel='application_type', ylabel='count'>
```



```
data.issue_d
```

```
0    Jan-15
1    Jan-15
2    Jan-15
```


```

3      Nov-14
4      Apr-13
...
396025  Oct-15
396026  Feb-15
396027  Oct-13
396028  Aug-12
396029  Jun-10
Name: issue_d, Length: 396030, dtype: object

sns.countplot(x='term', hue='loan_status', data=data)

<Axes: xlabel='term', ylabel='count'>

```




```

((data[data['loan_status']=='Charged Off'].groupby('sub_grade')
['loan_status'].count())/(data.groupby('sub_grade')
['sub_grade'].count())).plot(kind='bar')


<Axes: xlabel='sub_grade'>

```




```
plt.figure(figsize=(10,5))
sns.barplot(x='sub_grade', y='loan_amnt', hue='loan_status',
data=data,order=['A1','A2','A3','A4','A5','B1','B2','B3','B4','B5','C1',
', 'C2','C3','C4','C5','D1','D2','D3','D4','D5','E1','E2','E3','E4','E5',
', 'F1','F2','F3','F4','F5','G1','G2','G3','G4','G5'])

<Axes: xlabel='sub_grade', ylabel='loan_amnt'>
```



```
plt.figure(figsize=(10,5))
sns.barplot(x='sub_grade', y='int_rate', hue='loan_status',
data=data,order=['A1','A2','A3','A4','A5','B1','B2','B3','B4','B5','C1',
', 'C2','C3','C4','C5','D1','D2','D3','D4','D5','E1','E2','E3','E4','E5',
', 'F1','F2','F3','F4','F5','G1','G2','G3','G4','G5'])
```

<Axes: xlabel='sub_grade', ylabel='int_rate'>



#The 2 graphs above show that:

#a) the bank is putting more money in risky categories of borrowers

#b) the interest rates for fully paid borrowers and those charging off are same in each category, which means the innocent borrowers are being penalized for those who run away

#Recommendation:

#1. The bank should put more money into its group A & B customers

#2. For groups C & D where the default rate is high, the interest rate should be tailored according to the borrowers recent history even if it is not with the bank.

#This will be possible by collaborating with other banks and using a blockchain solution to identify behaviours that imply default, even though they are not financial related.

#So for example a customer's pattern of recent buying transactions in a retail store or in automobile purchases, might indicate that he will carry that same behaviour into his banking transactions.

```
data.isnull().sum()
```

loan_amnt	0
term	0
int_rate	0
installment	0
grade	0
sub_grade	0
emp_title	22927
emp_length	18301
home_ownership	0
annual_inc	0
verification_status	0
issue_d	0
loan_status	0
purpose	0
title	1756
dti	0
earliest_cr_line	0
open_acc	0
pub_rec	0
revol_bal	0
revol_util	276
total_acc	0
initial_list_status	0
application_type	0
mort_acc	37795
pub_rec_bankruptcies	535
address	0

dtype: int64

```
data.title.unique() #This is similar to the purpose column and can be deleted
```

```
array(['Vacation', 'Debt consolidation', 'Credit card
refinancing', ...,
      'Credit buster ', 'Loanforpayoff', 'Toxic Debt Payoff'],
      dtype=object)
```

```
data.title.value_counts()
```

```
title
Debt consolidation      152472
Credit card refinancing  51487
Home improvement        15264
Other                   12930
Debt Consolidation      11608
...
PayOffHighIntCreditCards    1
Heat my home                 1
Graduation/Travel Expenses  1
Daughter's Wedding Bill     1
Toxic Debt Payoff            1
Name: count, Length: 48804, dtype: int64
```

```
data.emp_title.unique() #This does not provide any ideas on loan
repayment, so I can delete
```


```
array(['Marketing', 'Credit analyst ', 'Statistician', ...,
      "Michael's Arts & Crafts", 'licensed bankere',
      'Gracon Services, Inc'], dtype=object)
```

```
data.emp_length.value_counts()
```

```
emp_length
10+ years    126041
2 years      35827
< 1 year     31725
3 years      31665
5 years      26495
1 year       25882
4 years      23952
6 years      20841
7 years      20819
8 years      19168
9 years      15314
Name: count, dtype: int64
```

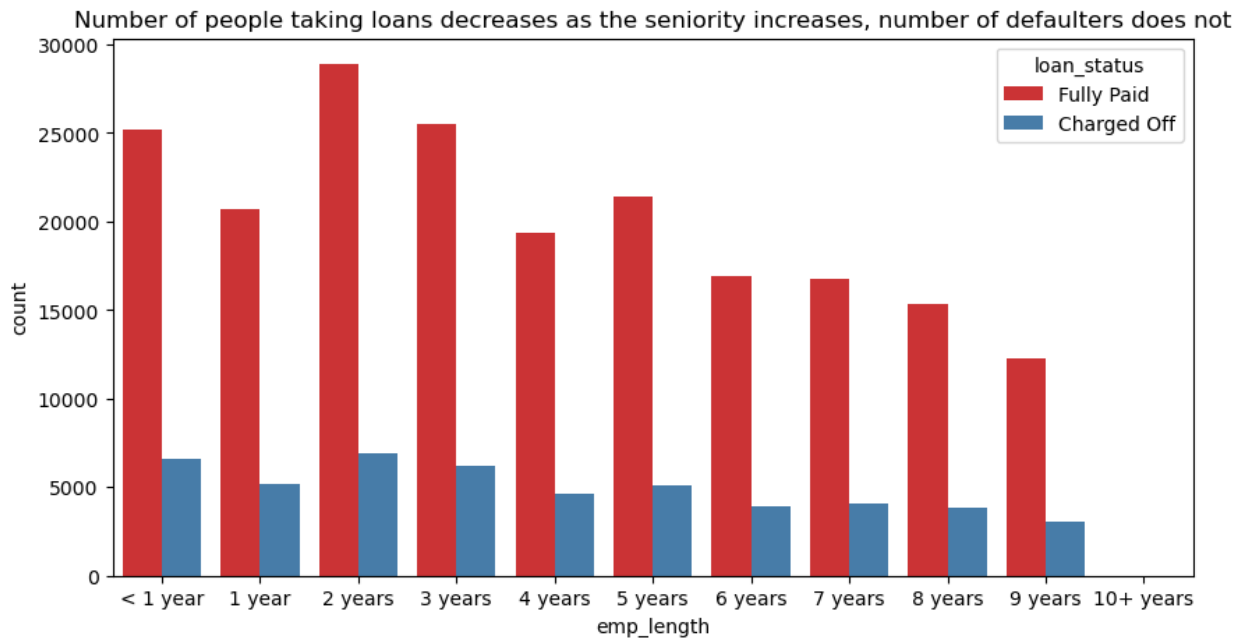
```
plt.figure(figsize=(10,5))
sns.countplot(hue='loan_status', x='emp_length', order=['< 1 year', '1
year', '2 years', '3 years', '4 years', '5 years', '6 years', '7
years', '8 years', '9 years', '10+ years'], data=data, palette='Set1')
```

<Axes: xlabel='emp_length', ylabel='count'>



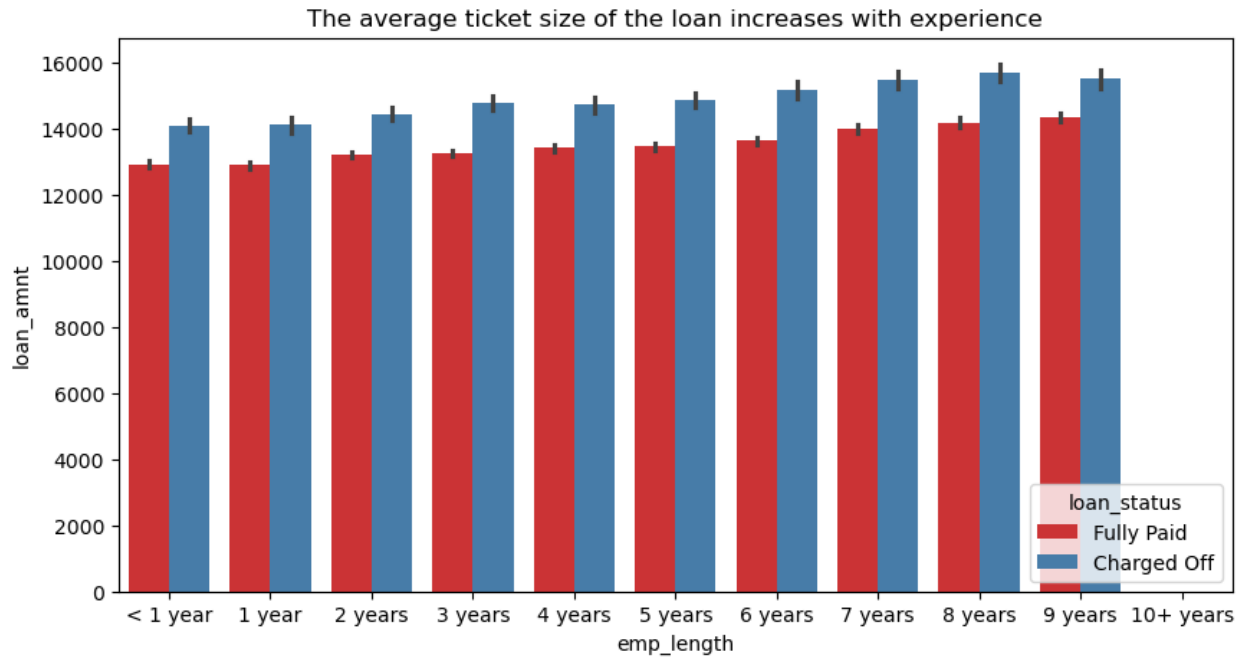
```
el_new=data[data['emp_length']!='10+ years']

plt.figure(figsize=(10,5))
sns.countplot(x='emp_length', hue='loan_status',order=['< 1 year', '1 year', '2 years','3 years', '4 years', '5 years', '6 years', '7 years', '8 years','9 years', '10+ years'],data=el_new, palette='Set1')
plt.title('Number of people taking loans decreases as the seniority increases, number of defaulters does not')
plt.show()
```




#The graph shows that as people spend more time in the company, the number of proportion of people paying the loan back is reducing, as long as it is <=10 years
#So my recommendation is to have a higher interest rate for people who are longer serving

```
plt.figure(figsize=(10,5))
sns.barplot(x='emp_length', y='loan_amnt', hue='loan_status',order=['< 1 year', '1 year', '2 years','3 years', '4 years', '5 years', '6 years', '7 years', '8 years','9 years', '10+ years'],data=el_new, palette='Set1')
plt.title('The average ticket size of the loan increases with experience')
plt.show()
```





```
sns.scatterplot(x='dti',y='loan_amnt', hue='loan_status', data=data)  
<Axes: xlabel='dti', ylabel='loan_amnt'>
```




```
data.dti.hist()
```

```
<Axes: >
```




```
sns.scatterplot(x='annual_inc', y='dti', hue='loan_status', data=data)
```

```
<Axes: xlabel='annual_inc', ylabel='dti'>
```




```
sns.barplot(x='loan_status', y='dti', data=data)  
<Axes: xlabel='loan_status', ylabel='dti'>
```




```
sns.barplot(x='loan_status',y='revol_bal', data=data) #This is too small a difference so the feature can be removed
```

```
<Axes: xlabel='loan_status', ylabel='revol_bal'>
```




```
sns.barplot(x='loan_status', y='revol_util', data=data) #This is a %  
figure so it makes sense to keep it, the 5% or so difference is  
significant
```


```
<Axes: xlabel='loan_status', ylabel='revol_util'>
```




```
sns.scatterplot(x='revol_util', y='loan_amnt', hue='loan_status',  
data=data) #There appears to be no correlation between the revolving  
utilization and the loan status or loan amount, so we can remove it  
<Axes: xlabel='revol_util', ylabel='loan_amnt'>
```



```
sns.barplot(x='pub_rec_bankruptcies', y='loan_amnt',  
hue='loan_status', data=data)  
#So in any case, the money loaned to people who defect is more than  
that loaned to people who repay  
#Also if the number of bankruptcies is above 5, the bank is giving  
even more money to the defaulters relative to the people who are  
paying  
<Axes: xlabel='pub_rec_bankruptcies', ylabel='loan_amnt'>
```




```
sns.barplot(x='pub_rec', y='loan_amnt', hue='loan_status', data=data)  
<Axes: xlabel='pub_rec', ylabel='loan_amnt'>
```





```
data.installment.hist()
```


```
<Axes: >
```



```
sns.barplot(x='loan_status', y='installment', data=data)  
<Axes: xlabel='loan_status', ylabel='installment'>
```



```
sns.barplot(hue='term', y='installment', x='loan_status', data=data)  
<Axes: xlabel='loan_status', ylabel='installment'>
```



```
features_to_drop=['emp_title', 'issue_d', 'loan_status',  
'title', 'earliest_cr_line', 'open_acc',  
'pub_rec_bankruptcies', 'pub_rec', 'revol_util', 'revol_bal',  
'total_acc', 'installment', 'mort_acc', 'address', 'grade']
```

```
data.emp_length.isnull().sum()/data.shape[0] #only 4.6% of the records  
are missing the employee length column, so I will drop the missing  
rows
```

```
0.046211145620281294
```

```
data.drop(data[data.emp_length.isnull()].index, inplace=True)
```

```
data=data[data.emp_length!='10+ years']
```

```
x=data.drop(features_to_drop, axis=1)
```

```
x.shape
```

```
(251688, 12)
```

```
x.isnull().sum()
```

```
loan_amnt      0
term           0
int_rate       0
sub_grade      0
emp_length     0
home_ownership 0
annual_inc     0
verification_status 0
purpose        0
dti            0
initial_list_status 0
application_type 0
dtype: int64
```

```
#x[x.emp_length.notnull()]['loan_amnt'].describe()
```

```
y=data['loan_status']
```

```
y
```

```
1      Fully Paid
2      Fully Paid
3      Fully Paid
4      Charged Off
6      Fully Paid
```

```
...
396017  Fully Paid
396022  Fully Paid
396024  Fully Paid
396025  Fully Paid
396026  Fully Paid
```

```
Name: loan_status, Length: 251688, dtype: object
```

```
y.replace(['Fully Paid','Charged Off'],[0,1], inplace=True)
```

```
/var/folders/c5/jzdz65t53qj6w2j6h9rj5sr80000gn/T/
ipykernel_8360/4151647945.py:1: FutureWarning: Downcasting behavior in
`replace` is deprecated and will be removed in a future version. To
retain the old behavior, explicitly call
`result.infer_objects(copy=False)`. To opt-in to the future behavior,
set `pd.set_option('future.no_silent_downcasting', True)`
  y.replace(['Fully Paid','Charged Off'],[0,1], inplace=True)
```

```
x.columns
```

```
Index(['loan_amnt', 'term', 'int_rate', 'sub_grade', 'emp_length',
      'home_ownership', 'annual_inc', 'verification_status',
      'purpose', 'dti',
```

```
'initial_list_status', 'application_type'],  
dtype='object')
```

```
x.dtypes
```

```
loan_amnt      int64  
term           object  
int_rate       float64  
sub_grade      object  
emp_length     object  
home_ownership object  
annual_inc     float64  
verification_status object  
purpose        object  
dti            float64  
initial_list_status object  
application_type object  
dtype: object
```


```
#Cleaning the data
```

```
x.loan_amnt.describe()
```


```
count    251688.000000  
mean      13703.230885  
std       8183.898226  
min        500.000000  
25%       7500.000000  
50%      12000.000000  
75%      18550.000000  
max      40000.000000  
Name: loan_amnt, dtype: float64
```

```
x.loan_amnt.hist()
```

```
<Axes: >
```



```
iqr_la=x.loan_amnt.quantile(.75)-x.loan_amnt.quantile(.25)
uql=x.loan_amnt.quantile(.75)+(1.5*iqr_la)
uql
35125.0
x.loan_amnt=[uql if i>uql else i for i in x.loan_amnt]
x.loan_amnt.hist()
<Axes: >
```



```
import scipy
from scipy.stats import shapiro

shapiro(x.loan_amnt) #This shows the data is not normal


/opt/anaconda3/lib/python3.12/site-packages/scipy/stats/_axis_nan_policy.py:531: UserWarning: scipy.stats.shapiro: For N >
5000, computed p-value may not be accurate. Current N is 251688.
  res = hypotest_fun_out(*samples, **kws)

ShapiroResult(statistic=0.9335602104856013,
pvalue=1.4959779862908053e-119)

x.loan_amnt=np.sqrt(x.loan_amnt)

x.loan_amnt.hist() #This looks normal now

<Axes: >
```





```
x.annual_inc.describe()
```


```
count    2.516880e+05
mean     7.246300e+04
std      5.657586e+04
min      4.000000e+03
25%      4.400000e+04
50%      6.000000e+04
75%      8.700000e+04
max      7.600000e+06
Name: annual_inc, dtype: float64
```

```
x.annual_inc.hist()
```

```
<Axes: >
```



```
ann_inc_iqr=x.annual_inc.quantile(.75)-x.annual_inc.quantile(.25)
uql_ann_inc=x.annual_inc.quantile(.75)+(1.5*ann_inc_iqr)
x.annual_inc=[uql_ann_inc if i>=uql_ann_inc else i for i in
x.annual_inc]
x.annual_inc=np.sqrt(x.annual_inc) #This looks more normal to me than
earlier
x.annual_inc.hist()
<Axes: >
```



```
sorted(x.sub_grade.unique())
```

```
['A1',  
'A2',  
'A3',  
'A4',  
'A5',  
'B1',  
'B2',  
'B3',  
'B4',  
'B5',  
'C1',  
'C2',  
'C3',  
'C4',  
'C5',  
'D1',  
'D2',  
'D3',  
'D4',  
'D5',  
'E1',  
'E2',  
'E3',
```

```
'E4',  
'E5',  
'F1',  
'F2',  
'F3',  
'F4',  
'F5',  
'G1',  
'G2',  
'G3',  
'G4',  
'G5']
```

```
x.sub_grade.replace(['A1','A2','A3','A4','A5','B1','B2','B3','B4','B5',  
, 'C1','C2','C3','C4','C5','D1','D2','D3','D4','D5','E1','E2','E3','E4',  
, 'E5','F1','F2','F3','F4','F5','G1','G2','G3','G4','G5'],list(range(35  
)), inplace=True)
```

/var/folders/c5/jzdz65t53qj6w2j6h9rj5sr80000gn/T/

ipykernel_8360/2111346709.py:1: FutureWarning: A value is trying to be set on a copy of a DataFrame or Series through chained assignment using an inplace method.

The behavior will change in pandas 3.0. This inplace method will never work because the intermediate object on which we are setting values always behaves as a copy.

For example, when doing 'df[col].method(value, inplace=True)', try using 'df.method({col: value}, inplace=True)' or df[col] = df[col].method(value) instead, to perform the operation inplace on the original object.

```
x.sub_grade.replace(['A1','A2','A3','A4','A5','B1','B2','B3','B4','B5',  
, 'C1','C2','C3','C4','C5','D1','D2','D3','D4','D5','E1','E2','E3','E4',  
, 'E5','F1','F2','F3','F4','F5','G1','G2','G3','G4','G5'],list(range(35  
)), inplace=True)
```

/var/folders/c5/jzdz65t53qj6w2j6h9rj5sr80000gn/T/ipykernel_8360/2111346709.py:1: FutureWarning: Downcasting behavior in 'replace' is deprecated and will be removed in a future version. To retain the old behavior, explicitly call 'result.infer_objects(copy=False)'. To opt-in to the future behavior, set 'pd.set_option('future.no_silent_downcasting', True)'

```
x.sub_grade.replace(['A1','A2','A3','A4','A5','B1','B2','B3','B4','B5',  
, 'C1','C2','C3','C4','C5','D1','D2','D3','D4','D5','E1','E2','E3','E4',  
, 'E5','F1','F2','F3','F4','F5','G1','G2','G3','G4','G5'],list(range(35  
)), inplace=True)
```

```
x.sub_grade
```

```

1          9
2          7
3          1
4         14
6          0
..
396017     7
396022    10
396024     8
396025     8
396026    10
Name: sub_grade, Length: 251688, dtype: int64

```

```
x.home_ownership.value_counts()
```

```

home_ownership
RENT          117718
MORTGAGE      111973
OWN           21893
OTHER           83
NONE           18
ANY             3
Name: count, dtype: int64

```

#I am clubbing the categories "ANY" with "OTHER" in home ownership as otherwise the categories will be too small

```
x.home_ownership.replace('ANY', 'OTHER', inplace=True)
```

/var/folders/c5/jzdz65t53qj6w2j6h9rj5sr80000gn/T/ipykernel_8360/336726384.py:2: FutureWarning: A value is trying to be set on a copy of a DataFrame or Series through chained assignment using an inplace method.
The behavior will change in pandas 3.0. This inplace method will never work because the intermediate object on which we are setting values always behaves as a copy.

For example, when doing 'df[col].method(value, inplace=True)', try using 'df.method({col: value}, inplace=True)' or df[col] = df[col].method(value) instead, to perform the operation inplace on the original object.

```
x.home_ownership.replace('ANY', 'OTHER', inplace=True)
```

```
x.home_ownership.value_counts()
```

```

home_ownership
RENT          117718
MORTGAGE      111973
OWN           21893
OTHER           86

```

```
NONE          18
Name: count, dtype: int64
```

```
dummies_ho=pd.get_dummies(x.home_ownership, drop_first=True)
x=pd.concat([x.drop('home_ownership', axis=1), dummies_ho], axis=1)
x.verification_status.unique()
#since there is an order here, source verified is better than not
verified and verified is better than source verified, I will replace
it as 0,1,2
```

```
array(['Not Verified', 'Source Verified', 'Verified'], dtype=object)
```

```
x.verification_status.replace(['Not Verified', 'Source Verified',
'Verified'],[0,1,2], inplace=True)
```


```
/var/folders/c5/jzdz65t53qj6w2j6h9rj5sr80000gn/T/
ipykernel_8360/1240797381.py:1: FutureWarning: A value is trying to be
set on a copy of a DataFrame or Series through chained assignment
using an inplace method.
The behavior will change in pandas 3.0. This inplace method will never
work because the intermediate object on which we are setting values
always behaves as a copy.
```

For example, when doing 'df[col].method(value, inplace=True)', try using 'df.method({col: value}, inplace=True)' or df[col] = df[col].method(value) instead, to perform the operation inplace on the original object.


```
    x.verification_status.replace(['Not Verified', 'Source Verified',
'Verified'],[0,1,2], inplace=True)
/var/folders/c5/jzdz65t53qj6w2j6h9rj5sr80000gn/T/ipykernel_8360/124079
7381.py:1: FutureWarning: Downcasting behavior in `replace` is
deprecated and will be removed in a future version. To retain the old
behavior, explicitly call `result.infer_objects(copy=False)`. To opt-
in to the future behavior, set
`pd.set_option('future.no_silent_downcasting', True)`
    x.verification_status.replace(['Not Verified', 'Source Verified',
'Verified'],[0,1,2], inplace=True)
```

```
dummies_purpose=pd.get_dummies(x.purpose, drop_first=True)
x=pd.concat([x.drop('purpose', axis=1), dummies_purpose], axis=1)
x.dti.hist()
```

```
<Axes: >
```



```
iqr_dti=x.dti.quantile(.75)-x.dti.quantile(.25)
uql_dti=x.dti.quantile(.75)+(1.5*iqr_dti)
x.dti=[uql_dti if i>=uql_dti else i for i in x.dti]
x.dti.hist()
<Axes: >
```



```
x.initial_list_status.replace(['w','f'],[0,1], inplace=True)
```

```
/var/folders/c5/jzdz65t53qj6w2j6h9rj5sr80000gn/T/
```

```
ipykernel_8360/1420288000.py:1: FutureWarning: A value is trying to be
set on a copy of a DataFrame or Series through chained assignment
using an inplace method.
```

The behavior will change in pandas 3.0. This inplace method will never work because the intermediate object on which we are setting values always behaves as a copy.

For example, when doing 'df[col].method(value, inplace=True)', try using 'df.method({col: value}, inplace=True)' or df[col] = df[col].method(value) instead, to perform the operation inplace on the original object.

```
x.initial_list_status.replace(['w','f'],[0,1], inplace=True)
/var/folders/c5/jzdz65t53qj6w2j6h9rj5sr80000gn/T/ipykernel_8360/142028
8000.py:1: FutureWarning: Downcasting behavior in `replace` is
deprecated and will be removed in a future version. To retain the old
behavior, explicitly call `result.infer_objects(copy=False)`. To opt-
in to the future behavior, set
`pd.set_option('future.no_silent_downcasting', True)`
x.initial_list_status.replace(['w','f'],[0,1], inplace=True)
x.initial_list_status.unique()
```



```

array([1, 0])

x.application_type.value_counts()

application_type
INDIVIDUAL      251285
JOINT            226
DIRECT_PAY       177
Name: count, dtype: int64

dummies_app_type=pd.get_dummies(x.application_type, drop_first=True)

x=pd.concat([x.drop('application_type', axis=1), dummies_app_type],
axis=1)

x.emp_length.unique()

array(['4 years', '< 1 year', '6 years', '9 years', '2 years', '3
years',
      '8 years', '7 years', '5 years', '1 year'], dtype=object)

x.emp_length.value_counts()

emp_length
2 years      35827
< 1 year     31725
3 years      31665
5 years      26495
1 year       25882
4 years      23952
6 years      20841
7 years      20819
8 years      19168
9 years      15314
Name: count, dtype: int64

x.replace('< 1 year',"0", inplace=True)

x.emp_length=x.emp_length.str.split(expand=True)[0].astype('int64')

x.emp_length


1          4
2          0
3          6
4          9
6          2
..
396017     8
396022     1
396024     5
396025     2

```

```
396026    5
Name: emp_length, Length: 251688, dtype: int64


x.emp_length.hist()

<Axes: >
```



```
iqr_el=x.emp_length.quantile(.75)-x.emp_length.quantile(.25)
ul=x.emp_length.quantile(.75)+(1.5*iqr_el)
ll=x.emp_length.quantile(.25)-(1.5*iqr_el)
ul
12.0
ll
-4.0
x.emp_length=[ul if i>=ul else i for i in x.emp_length]
x.emp_length.hist() #This still does not look normal enough and I have
let go after trying multiple transforms

<Axes: >
```



```
import sklearn
from sklearn.preprocessing import StandardScaler
from sklearn.model_selection import train_test_split
```

```
import statsmodels
from statsmodels.stats.outliers_influence import
variance_inflation_factor as vif
```

```
x.shape
```

```
(251688, 28)
```

```
x.term=x.term.str.split(expand=True)[0].replace(['36','60'],
[0,1]).astype('int64')
```

```
/var/folders/c5/jzdz65t53qj6w2j6h9rj5sr80000gn/T/
```


```
ipykernel_8360/2232850954.py:1: FutureWarning: Downcasting behavior in
`replace` is deprecated and will be removed in a future version. To
retain the old behavior, explicitly call
`result.infer_objects(copy=False)`. To opt-in to the future behavior,
set `pd.set_option('future.no_silent_downcasting', True)`
```

```
x.term=x.term.str.split(expand=True)[0].replace(['36','60'],  
[0,1]).astype('int64')
```

```
x.replace('other','OTHER',inplace=True)
```

```
x.term.hist()
```

<Axes: >



x.dtypes

loan_amnt	float64
term	int64
int_rate	float64
sub_grade	int64
emp_length	int64
annual_inc	float64
verification_status	int64
dti	float64
initial_list_status	int64
NONE	bool
OTHER	bool
OWN	bool
RENT	bool
credit_card	bool

```

debt_consolidation    bool
educational            bool
home_improvement      bool
house                 bool
major_purchase        bool
medical               bool
moving                bool
other                 bool
renewable_energy      bool
small_business         bool
vacation              bool
wedding               bool
INDIVIDUAL            bool
JOINT                 bool
dtype: object

```

```
x.columns
```

```

Index(['loan_amnt', 'term', 'int_rate', 'sub_grade', 'emp_length',
       'annual_inc', 'verification_status', 'dti',
       'initial_list_status',
       'NONE', 'OTHER', 'OWN', 'RENT', 'credit_card',
       'debt_consolidation',
       'educational', 'home_improvement', 'house', 'major_purchase',
       'medical',
       'moving', 'other', 'renewable_energy', 'small_business',
       'vacation',
       'wedding', 'INDIVIDUAL', 'JOINT'],
      dtype='object')

```

```

x[['loan_amnt', 'term', 'int_rate', 'sub_grade', 'emp_length',
    'annual_inc', 'verification_status', 'dti',
    'initial_list_status',
    'NONE', 'OTHER', 'OWN', 'RENT', 'credit_card',
    'debt_consolidation',
    'educational', 'home_improvement', 'house', 'major_purchase',
    'medical',
    'moving', 'other', 'renewable_energy', 'small_business',
    'vacation',
    'wedding', 'INDIVIDUAL', 'JOINT']] = x[['loan_amnt', 'term',
    'int_rate', 'sub_grade', 'emp_length',
    'annual_inc', 'verification_status', 'dti',
    'initial_list_status',
    'NONE', 'OTHER', 'OWN', 'RENT', 'credit_card',
    'debt_consolidation',
    'educational', 'home_improvement', 'house', 'major_purchase',
    'medical',
    'moving', 'other', 'renewable_energy', 'small_business',
    'vacation',
    'wedding', 'INDIVIDUAL', 'JOINT']].astype('int64')

```

```
x.dtypes
```

```
loan_amnt      int64
term           int64
int_rate       int64
sub_grade      int64
emp_length     int64
annual_inc     int64
verification_status int64
dti            int64
initial_list_status int64
NONE           int64
OTHER          int64
OWN            int64
RENT           int64
credit_card    int64
debt_consolidation int64
educational    int64
home_improvement int64
house          int64
major_purchase int64
medical        int64
moving         int64
other          int64
renewable_energy int64
small_business int64
vacation       int64
wedding        int64
INDIVIDUAL     int64
JOINT          int64
dtype: object
```

```
#Checking for multicollinearity
```

```
vif_df=pd.DataFrame()
```

```
vif_df['Features']=x.columns
```

```
vif_df['vif']=[vif(x.values,i) for i in range(x.shape[1])]
```

```
vif_df.sort_values(by='vif',ascending=False)
```

	Features	vif
2	int_rate	159.118540
26	INDIVIDUAL	122.039449
3	sub_grade	66.343499
14	debt_consolidation	43.915574
5	annual_inc	27.842232
0	loan_amnt	20.483719
13	credit_card	16.526390
7	dti	5.957414

16	home_improvement	5.055984
21	other	4.895489
4	emp_length	3.097130
6	verification_status	2.869599
8	initial_list_status	2.749938
18	major_purchase	2.736055
12	RENT	2.347416
23	small_business	2.227005
1	term	1.964752
19	medical	1.757477
20	moving	1.636497
17	house	1.459496
25	wedding	1.437254
24	vacation	1.432592
11	OWN	1.222579
27	JOINT	1.108100
22	renewable_energy	1.064330
15	educational	1.063408
10	OTHER	1.001546
9	NONE	1.000305

```
x_new=x.drop('int_rate', axis=1)
```

```
vif_df2=pd.DataFrame()
```

```
vif_df2['features']=x_new.columns
```

```
vif_df2['vif']=[vif(x_new.values,i) for i in range(x_new.shape[1])]
```

```
vif_df2.sort_values(by='vif', ascending=False)
```

	features	vif
25	INDIVIDUAL	100.644484
13	debt_consolidation	43.814229
4	annual_inc	27.804375
0	loan_amnt	20.478800
12	credit_card	16.479834
6	dti	5.955760
2	sub_grade	5.552463
15	home_improvement	5.049818
20	other	4.890019
3	emp_length	3.089352
5	verification_status	2.862165
17	major_purchase	2.734204
7	initial_list_status	2.673364
11	RENT	2.347366
22	small_business	2.226757
1	term	1.948089
18	medical	1.756585
19	moving	1.635659
16	house	1.458538

24	wedding	1.436356
23	vacation	1.431907
10	OWN	1.222479
26	JOINT	1.090050
21	renewable_energy	1.064305
14	educational	1.063318
9	OTHER	1.001534
8	NONE	1.000261

```
vif_df3=pd.DataFrame()
```

```
x_new_2=x_new.drop('INDIVIDUAL', axis=1)
```

```
vif_df3['features']=x_new_2.columns
```


```
vif_df3['vif']=[vif(x_new_2.values, i) for i in range(x_new_2.shape[1])]
```

```
vif_df3.sort_values(by='vif', ascending=False)
```

	features	vif
4	annual_inc	23.857701
0	loan_amnt	20.447269
13	debt_consolidation	17.852452
12	credit_card	6.946988
6	dti	5.693950
2	sub_grade	5.491294
3	emp_length	3.048729
5	verification_status	2.859252
7	initial_list_status	2.593314
15	home_improvement	2.520120
20	other	2.467204
11	RENT	2.234134
1	term	1.942312
17	major_purchase	1.620227
22	small_business	1.506574
18	medical	1.288669
19	moving	1.251106
10	OWN	1.204696
16	house	1.188007
24	wedding	1.168347
23	vacation	1.166097
21	renewable_energy	1.024432
14	educational	1.019822
25	JOINT	1.003133
9	OTHER	1.001444
8	NONE	1.000256

```
sns.scatterplot(x=x.annual_inc,y=x.loan_amnt,hue=y, data=data)
```

```
<Axes: xlabel='annual_inc', ylabel='loan_amnt'>
```

```
np.corrcoef(x.annual_inc, x.loan_amnt)
array([[1.          , 0.51401461],
       [0.51401461, 1.          ]])

x_new_3=x_new_2.drop('annual_inc', axis=1)
vif_df_4=pd.DataFrame(columns=['features','vif'])
vif_df_4['features']=x_new_3.columns
vif_df_4['vif']=[vif(x_new_3.values, i) for i in
range(x_new_3.shape[1])]

vif_df_4.sort_values(by='vif', ascending=False) #I will keep the loan
amount column as it is an important variable to determines who pays up
na who runs away
```

	features	vif
0	loan_amnt	14.192121
12	debt_consolidation	13.206193
5	dti	5.518560
2	sub_grade	5.465802
11	credit_card	5.234918
3	emp_length	3.030336
4	verification_status	2.858417

6	initial_list_status	2.592482
10	RENT	2.212688
14	home_improvement	1.968034
19	other	1.942928
1	term	1.932355
16	major_purchase	1.389450
21	small_business	1.361115
9	OWN	1.199809
17	medical	1.176443
18	moving	1.152466
15	house	1.134461
23	wedding	1.114384
22	vacation	1.095230
20	renewable_energy	1.015743
13	educational	1.013799
24	JOINT	1.002837
8	OTHER	1.001436
7	NONE	1.000245

```
y.value_counts() #This is a little imbalanced
```

```
loan_status
0    202268
1     49420
Name: count, dtype: int64
```

```
import statsmodels.api as sm
from sklearn.linear_model import LogisticRegression

xtrain, xtest, ytrain, ytest=train_test_split(x_new_3, y,
test_size=0.2, random_state=10)

sc=StandardScaler()

xtrain_sc=sc.fit_transform(xtrain)

xtest_sc=sc.transform(xtest)

xtrain_sc=pd.DataFrame(xtrain_sc, columns=xtrain.columns)

xtrain_sc
```

	loan_amnt	term	sub_grade	emp_length	
verification_status \					
0	-1.382222	-0.539094	-0.920924	1.128624	-
1.200334					
1	-0.064616	-0.539094	-0.768967	-0.692137	
1.264304					
2	-1.181716	-0.539094	-0.313096	-0.327985	-

1.200334					
3	-0.064616	-0.539094	-0.009182	0.036167	
0.031985					
4	-0.780706	-0.539094	-1.376794	1.492777	-
1.200334					
...
.					
201345	-0.809350	-0.539094	0.750602	1.492777	
0.031985					
201346	-1.181716	-0.539094	-0.313096	1.856929	
0.031985					
201347	-0.064616	-0.539094	-1.376794	-0.327985	
1.264304					
201348	-0.465626	-0.539094	0.294731	-0.692137	-
1.200334					
201349	2.169585	-0.539094	-0.768967	0.400320	
1.264304					

	dti	initial_list_status	NONE	OTHER	
OWN	...	\			
0	-1.313105		-1.264769	-0.008631	-0.019174
3.244807	...				
1	-0.073361		0.790658	-0.008631	-0.019174 -
0.308185	...				
2	-1.189130		-1.264769	-0.008631	-0.019174 -
0.308185	...				
3	-0.445284		0.790658	-0.008631	-0.019174 -
0.308185	...				
4	0.546511		0.790658	-0.008631	-0.019174 -
0.308185	...				
...
..					
201345	0.794460		0.790658	-0.008631	-0.019174 -
0.308185	...				
201346	0.546511		0.790658	-0.008631	-0.019174 -
0.308185	...				
201347	-1.313105		0.790658	-0.008631	-0.019174 -
0.308185	...				
201348	1.290358		0.790658	-0.008631	-0.019174 -
0.308185	...				
201349	-0.197335		-1.264769	-0.008631	-0.019174 -
0.308185	...				

	house	major_purchase	medical	moving	other	\
0	-0.078686	-0.157221	-0.102539	-0.092904	-0.237850	
1	-0.078686	-0.157221	-0.102539	-0.092904	-0.237850	
2	-0.078686	-0.157221	-0.102539	-0.092904	-0.237850	
3	-0.078686	-0.157221	-0.102539	-0.092904	-0.237850	
4	-0.078686	-0.157221	-0.102539	-0.092904	4.204328	

```

...
201345 -0.078686 -0.157221 -0.102539 -0.092904 -0.237850
201346 -0.078686 -0.157221 -0.102539 -0.092904 -0.237850
201347 -0.078686 -0.157221 -0.102539 -0.092904 -0.237850
201348 -0.078686 -0.157221 -0.102539 -0.092904 -0.237850
201349 -0.078686 -0.157221 -0.102539 -0.092904 -0.237850

```

```

renewable_energy small_business vacation wedding
JOINT
0 -0.029578 -0.130842 -0.077008 -0.077722 -
0.029662
1 -0.029578 -0.130842 -0.077008 -0.077722 -
0.029662
2 -0.029578 -0.130842 -0.077008 -0.077722 -
0.029662
3 -0.029578 -0.130842 -0.077008 -0.077722 -
0.029662
4 -0.029578 -0.130842 -0.077008 -0.077722 -
0.029662
...
.
201345 -0.029578 -0.130842 -0.077008 -0.077722 -
0.029662
201346 -0.029578 -0.130842 -0.077008 12.866336 -
0.029662
201347 -0.029578 -0.130842 -0.077008 -0.077722 -
0.029662
201348 -0.029578 -0.130842 -0.077008 -0.077722 -
0.029662
201349 -0.029578 -0.130842 -0.077008 -0.077722 -
0.029662

```

[201350 rows x 25 columns]

```
xtest_sc=pd.DataFrame(xtest_sc, columns=xtest.columns)
```

```
xtest_sc
```

```

loan_amnt term sub_grade emp_length verification_status
\
0 2.169585 -0.539094 1.054516 -0.692137 1.264304
1 0.851979 1.854964 -0.465053 1.128624 1.264304
2 0.307751 -0.539094 -0.161139 1.856929 1.264304
3 -0.809350 -0.539094 2.726042 0.400320 1.264304
4 -1.181716 -0.539094 -0.617010 1.856929 0.031985
...

```

50333	-0.866637	-0.539094	0.142775	0.764472	-1.200334
50334	-0.064616	1.854964	0.446688	0.036167	1.264304
50335	-0.809350	-0.539094	0.598645	-1.420442	-1.200334
50336	-0.236478	1.854964	0.142775	0.764472	1.264304
50337	-0.236478	-0.539094	-1.224837	-0.692137	1.264304

	dti	initial_list_status	NONE	OTHER
OWN ... \				
0	-0.197335	0.790658	-0.008631	-0.019174 -
0.308185 ...				
1	0.670486	0.790658	-0.008631	-0.019174 -
0.308185 ...				
2	1.786256	-1.264769	-0.008631	-0.019174 -
0.308185 ...				
3	0.794460	-1.264769	-0.008631	-0.019174 -
0.308185 ...				
4	1.786256	-1.264769	-0.008631	-0.019174 -
0.308185 ...				
...
...
50333	0.174588	0.790658	-0.008631	-0.019174 -
0.308185 ...				
50334	-0.197335	-1.264769	-0.008631	-0.019174 -
0.308185 ...				
50335	-1.685028	0.790658	-0.008631	-0.019174 -
0.308185 ...				
50336	-0.941182	-1.264769	-0.008631	-0.019174
3.244807 ...				
50337	-0.197335	0.790658	-0.008631	-0.019174 -
0.308185 ...				

	house	major_purchase	medical	moving	other	\
0	-0.078686	-0.157221	-0.102539	-0.092904	-0.23785	
1	12.708662	-0.157221	-0.102539	-0.092904	-0.23785	
2	-0.078686	-0.157221	-0.102539	-0.092904	-0.23785	
3	-0.078686	-0.157221	-0.102539	-0.092904	-0.23785	
4	-0.078686	-0.157221	-0.102539	-0.092904	-0.23785	
...	
50333	-0.078686	-0.157221	-0.102539	-0.092904	-0.23785	
50334	-0.078686	-0.157221	-0.102539	-0.092904	-0.23785	
50335	-0.078686	-0.157221	-0.102539	-0.092904	-0.23785	
50336	-0.078686	-0.157221	-0.102539	-0.092904	-0.23785	
50337	-0.078686	-0.157221	-0.102539	-0.092904	-0.23785	

	renewable_energy	small_business	vacation	wedding	JOINT
0	-0.029578	-0.130842	-0.077008	-0.077722	-0.029662
1	-0.029578	-0.130842	-0.077008	-0.077722	-0.029662
2	-0.029578	-0.130842	-0.077008	-0.077722	-0.029662
3	-0.029578	-0.130842	-0.077008	-0.077722	-0.029662
4	-0.029578	-0.130842	-0.077008	-0.077722	-0.029662
...
50333	-0.029578	-0.130842	-0.077008	-0.077722	-0.029662
50334	-0.029578	-0.130842	-0.077008	-0.077722	-0.029662
50335	-0.029578	-0.130842	-0.077008	-0.077722	-0.029662
50336	-0.029578	-0.130842	-0.077008	-0.077722	-0.029662
50337	-0.029578	-0.130842	-0.077008	-0.077722	-0.029662

```
[50338 rows x 25 columns]
```

```
xtrain_sc.shape
```

```
(201350, 25)
```

```
ytrain.shape
```

```
(201350,)
```

```
xtrain_sc=sm.add_constant(xtrain_sc)
```

```
model_logit=sm.Logit(list(ytrain), xtrain_sc).fit()
```

```
Optimization terminated successfully.
```

```
Current function value: 0.453414
```

```
Iterations 6
```

```
model_logit.summary()
```

```
<class 'statsmodels.iolib.summary.Summary'>
```

```
"""
```

```
Logit Regression Results
```

```
=====
=====
```

Dep. Variable: y No. Observations: 201350
 Model: Logit Df Residuals: 201324
 Method: MLE Df Model: 25
 Date: Fri, 18 Apr 2025 Pseudo R-squ.: 0.08555
 Time: 19:10:22 Log-Likelihood: -91295.
 converged: True LL-Null: -99836.
 Covariance Type: nonrobust LLR p-value: 0.000

		coef	std err	z	P> z	
[0.025 0.975]						

const		-1.5605	0.006	-247.090	0.000	-
1.573	-1.548					
loan_amnt		0.0279	0.007	4.031	0.000	
0.014	0.042					
term		0.1544	0.007	23.375	0.000	
0.141	0.167					
sub_grade		0.5340	0.007	77.854	0.000	
0.521	0.547					
emp_length		-0.0076	0.006	-1.272	0.203	-
0.019	0.004					
verification_status		-0.0087	0.006	-1.351	0.177	-
0.021	0.004					
dti		0.2370	0.006	39.622	0.000	
0.225	0.249					
initial_list_status		-0.0161	0.006	-2.700	0.007	-
0.028	-0.004					
NONE		0.0057	0.005	1.078	0.281	-
0.005	0.016					
OTHER		-0.0064	0.007	-0.871	0.384	-
0.021	0.008					
OWN		0.0577	0.006	9.422	0.000	
0.046	0.070					
RENT		0.1622	0.007	24.797	0.000	
0.149	0.175					
credit_card		0.0420	0.025	1.713	0.087	-
0.006	0.090					
debt_consolidation		0.0906	0.029	3.117	0.002	
0.034	0.148					
educational		0.0101	0.006	1.582	0.114	-

0.002	0.023				
home_improvement		0.0424	0.015	2.888	0.004
0.014	0.071				
house		-0.0087	0.007	-1.154	0.248
0.023	0.006				
major_purchase		0.0165	0.011	1.512	0.131
0.005	0.038				
medical		0.0184	0.008	2.254	0.024
0.002	0.034				
moving		0.0176	0.008	2.284	0.022
0.002	0.033				
other		0.0248	0.014	1.740	0.082
0.003	0.053				
renewable_energy		0.0136	0.006	2.430	0.015
0.003	0.025				
small_business		0.0688	0.009	7.517	0.000
0.051	0.087				
vacation		0.0071	0.007	0.960	0.337
0.007	0.021				
wedding		-0.0391	0.009	-4.555	0.000
0.056	-0.022				
JOINT		-0.0335	0.007	-4.780	0.000
0.047	-0.020				

```
=====
=====
"""
```

```
logit_summary=pd.DataFrame(model_logit.summary().tables[1]) #the
tables[0] and tables[1] give us the summary as a df
```

```
logit_summary.columns=logit_summary.iloc[0]
```

```
logit_summary
```

0		coef	std err	z	P> z
[0.025 \					
0		coef	std err	z	P> z
[0.025					
1	const	-1.5605	0.006	-247.090	0.000
-1.573					
2	loan_amnt	0.0279	0.007	4.031	0.000
0.014					
3	term	0.1544	0.007	23.375	0.000
0.141					
4	sub_grade	0.5340	0.007	77.854	0.000
0.521					
5	emp_length	-0.0076	0.006	-1.272	0.203
-0.019					

6	verification_status	-0.0087	0.006	-1.351	0.177
-0.021					
7	dti	0.2370	0.006	39.622	0.000
0.225					
8	initial_list_status	-0.0161	0.006	-2.700	0.007
-0.028					
9	NONE	0.0057	0.005	1.078	0.281
-0.005					
10	OTHER	-0.0064	0.007	-0.871	0.384
-0.021					
11	OWN	0.0577	0.006	9.422	0.000
0.046					
12	RENT	0.1622	0.007	24.797	0.000
0.149					
13	credit_card	0.0420	0.025	1.713	0.087
-0.006					
14	debt_consolidation	0.0906	0.029	3.117	0.002
0.034					
15	educational	0.0101	0.006	1.582	0.114
-0.002					
16	home_improvement	0.0424	0.015	2.888	0.004
0.014					
17	house	-0.0087	0.007	-1.154	0.248
-0.023					
18	major_purchase	0.0165	0.011	1.512	0.131
-0.005					
19	medical	0.0184	0.008	2.254	0.024
0.002					
20	moving	0.0176	0.008	2.284	0.022
0.002					
21	other	0.0248	0.014	1.740	0.082
-0.003					
22	renewable_energy	0.0136	0.006	2.430	0.015
0.003					
23	small_business	0.0688	0.009	7.517	0.000
0.051					
24	vacation	0.0071	0.007	0.960	0.337
-0.007					
25	wedding	-0.0391	0.009	-4.555	0.000
-0.056					
26	JOINT	-0.0335	0.007	-4.780	0.000
-0.047					
0	0.975]				
0	0.975]				
1	-1.548				
2	0.042				
3	0.167				
4	0.547				

5	0.004
6	0.004
7	0.249
8	-0.004
9	0.016
10	0.008
11	0.070
12	0.175
13	0.090
14	0.148
15	0.023
16	0.071
17	0.006
18	0.038
19	0.034
20	0.033
21	0.053
22	0.025
23	0.087
24	0.021
25	-0.022
26	-0.020

```
logit_summary=pd.DataFrame(model_logit.summary().tables[1]) #the values in this are not of float data type, they are of "cell" type
```

```
logit_summary.columns=logit_summary.iloc[0]
logit_summary.drop(0, inplace=True)
```

```
logit_summary.to_csv('logit_summary.csv')
#this is due to some different data type "cell" of the statsmodels table, was unable to convert to float directly
#so I am saving it as csv and reading it again
```

```
logit_summary=pd.read_csv('logit_summary.csv', index_col=0)
```

```
logit_summary.columns=['feature','coef','std err','z','pval',
'lower_crit_val', 'upper_crit_val']
```

```
logit_summary[logit_summary.pval>0.05] #getting those features where we are not confident because the p value shows >0.05
```

	feature	coef	std err	z	pval	lower_crit_val
5	emp_length	-0.0076	0.006	-1.272	0.203	-0.019
6	verification_status	-0.0087	0.006	-1.351	0.177	-0.021
9	NONE	0.0057	0.005	1.078	0.281	-0.005
10	OTHER	-0.0064	0.007	-0.871	0.384	-0.021

13	credit_card	0.0420	0.025	1.713	0.087	-0.006
15	educational	0.0101	0.006	1.582	0.114	-0.002
17	house	-0.0087	0.007	-1.154	0.248	-0.023
18	major_purchase	0.0165	0.011	1.512	0.131	-0.005
21	other	0.0248	0.014	1.740	0.082	-0.003
24	vacation	0.0071	0.007	0.960	0.337	-0.007

	upper_crit_val
5	0.004
6	0.004
9	0.016
10	0.008
13	0.090
15	0.023
17	0.006
18	0.038
21	0.053
24	0.021

#Here it shows that the OTHER feature has a p value that is very high so I will remove this and repeat

```
x_new_4=x_new_3.drop('OTHER', axis=1)
vif_df_5=pd.DataFrame(columns=['features','vif'])
vif_df_5.features=x_new_4.columns
vif_df_5.vif=[vif(x_new_4.values, i) for i in range(x_new_4.shape[1])]
vif_df_5.sort_values(by='vif', ascending=False)
```

	features	vif
0	loan_amnt	14.192116
11	debt_consolidation	13.205007
5	dti	5.518438
2	sub_grade	5.465608
10	credit_card	5.234321
3	emp_length	3.030180
4	verification_status	2.858404
6	initial_list_status	2.592180
9	RENT	2.211596
13	home_improvement	1.967961
18	other	1.942700
1	term	1.932223

15	major_purchase	1.389370
20	small_business	1.360946
8	OWN	1.199647
16	medical	1.176360
17	moving	1.152456
14	house	1.134435
22	wedding	1.114360
21	vacation	1.095225
19	renewable_energy	1.015680
12	educational	1.013569
23	JOINT	1.002837
7	NONE	1.000245

```
xtrain_new, xtest_new, ytrain, ytest=train_test_split(x_new_4, y,
test_size=0.2, random_state=10)
```

```
xtrain_sc_new=sc.fit_transform(xtrain_new)
```

```
xtest_sc_new=sc.transform(xtest_new)
```

```
xtrain_sc_new=pd.DataFrame(xtrain_sc_new, columns=xtrain_new.columns)
```

```
xtest_sc_new=pd.DataFrame(xtest_sc_new, columns=xtest_new.columns)
```

```
xtrain_sc_new=sm.add_constant(xtrain_sc_new)
```

```
model_logit_2=sm.Logit(list(ytrain), xtrain_sc_new).fit()
```

```
Optimization terminated successfully.
```

```
Current function value: 0.453416
```

```
Iterations 6
```

```
model_logit_2.summary()
```

```
<class 'statsmodels.iolib.summary.Summary'>
```

```
"""
```

Logit Regression Results

```
=====
```

```
=====
```

```
Dep. Variable:                y    No. Observations:
```

```
201350
```

```
Model:                Logit    Df Residuals:
```

```
201325
```

```
Method:                MLE    Df Model:
```

```
24
```

```
Date:                Fri, 18 Apr 2025    Pseudo R-squ.:
```

```
0.08554
```

```
Time:                19:10:47    Log-Likelihood:
```

```
-91295.
```

```
converged:                True    LL-Null:
```

-99836.

Covariance Type: nonrobust LLR p-value:

0.000

[0.025 0.975]		coef	std err	z	P> z	
const		-1.5604	0.006	-247.092	0.000	-
1.573	-1.548					
loan_amnt		0.0279	0.007	4.031	0.000	
0.014	0.042					
term		0.1545	0.007	23.385	0.000	
0.142	0.167					
sub_grade		0.5339	0.007	77.851	0.000	
0.520	0.547					
emp_length		-0.0075	0.006	-1.265	0.206	-
0.019	0.004					
verification_status		-0.0087	0.006	-1.350	0.177	-
0.021	0.004					
dti		0.2370	0.006	39.629	0.000	
0.225	0.249					
initial_list_status		-0.0161	0.006	-2.707	0.007	-
0.028	-0.004					
NONE		0.0057	0.005	1.079	0.281	-
0.005	0.016					
OWN		0.0578	0.006	9.434	0.000	
0.046	0.070					
RENT		0.1624	0.007	24.823	0.000	
0.150	0.175					
credit_card		0.0420	0.025	1.713	0.087	-
0.006	0.090					
debt_consolidation		0.0906	0.029	3.118	0.002	
0.034	0.148					
educational		0.0101	0.006	1.575	0.115	-
0.002	0.023					
home_improvement		0.0425	0.015	2.890	0.004	
0.014	0.071					
house		-0.0086	0.007	-1.154	0.249	-
0.023	0.006					
major_purchase		0.0165	0.011	1.513	0.130	-
0.005	0.038					
medical		0.0184	0.008	2.249	0.024	
0.002	0.034					
moving		0.0176	0.008	2.285	0.022	
0.003	0.033					
other		0.0248	0.014	1.740	0.082	-
0.003	0.053					

renewable_energy	0.0136	0.006	2.424	0.015	
0.003	0.025				
small_business	0.0688	0.009	7.514	0.000	
0.051	0.087				
vacation	0.0071	0.007	0.962	0.336	-
0.007	0.022				
wedding	-0.0391	0.009	-4.554	0.000	-
0.056	-0.022				
JOINT	-0.0335	0.007	-4.780	0.000	-
0.047	-0.020				

```
=====
=====
"""
```

```
model_logit_2.summary().tables[1]
<class 'statsmodels.iolib.table.SimpleTable'>

logit_summary_2=pd.DataFrame(model_logit_2.summary().tables[1])

logit_summary_2.columns=logit_summary.columns
logit_summary_2.drop(0, inplace=True)

logit_summary_2
```

	feature	coef	std err	z	pval	\
1	const	-1.5604	0.006	-247.092	0.000	
2	loan_amnt	0.0279	0.007	4.031	0.000	
3	term	0.1545	0.007	23.385	0.000	
4	sub_grade	0.5339	0.007	77.851	0.000	
5	emp_length	-0.0075	0.006	-1.265	0.206	
6	verification_status	-0.0087	0.006	-1.350	0.177	
7	dti	0.2370	0.006	39.629	0.000	
8	initial_list_status	-0.0161	0.006	-2.707	0.007	
9	NONE	0.0057	0.005	1.079	0.281	
10	OWN	0.0578	0.006	9.434	0.000	
11	RENT	0.1624	0.007	24.823	0.000	
12	credit_card	0.0420	0.025	1.713	0.087	
13	debt_consolidation	0.0906	0.029	3.118	0.002	
14	educational	0.0101	0.006	1.575	0.115	
15	home_improvement	0.0425	0.015	2.890	0.004	
16	house	-0.0086	0.007	-1.154	0.249	
17	major_purchase	0.0165	0.011	1.513	0.130	
18	medical	0.0184	0.008	2.249	0.024	
19	moving	0.0176	0.008	2.285	0.022	
20	other	0.0248	0.014	1.740	0.082	
21	renewable_energy	0.0136	0.006	2.424	0.015	
22	small_business	0.0688	0.009	7.514	0.000	
23	vacation	0.0071	0.007	0.962	0.336	
24	wedding	-0.0391	0.009	-4.554	0.000	
25	JOINT	-0.0335	0.007	-4.780	0.000	

	lower_crit_val	upper_crit_val
1	-1.573	-1.548
2	0.014	0.042
3	0.142	0.167
4	0.520	0.547
5	-0.019	0.004
6	-0.021	0.004
7	0.225	0.249
8	-0.028	-0.004
9	-0.005	0.016
10	0.046	0.070
11	0.150	0.175
12	-0.006	0.090
13	0.034	0.148
14	-0.002	0.023
15	0.014	0.071
16	-0.023	0.006
17	-0.005	0.038
18	0.002	0.034
19	0.003	0.033
20	-0.003	0.053
21	0.003	0.025
22	0.051	0.087
23	-0.007	0.022
24	-0.056	-0.022
25	-0.047	-0.020

```
logit_summary_2.to_csv('logit_summary_2.csv')
```

```
logit_summary_2=pd.read_csv('logit_summary_2.csv', index_col=0)
```

```
logit_summary_2
```

	feature	coef	std err	z	pval	
lower_crit_val	\					
1	const	-1.5604	0.006	-247.092	0.000	-
1.573						
2	loan_amnt	0.0279	0.007	4.031	0.000	
0.014						
3	term	0.1545	0.007	23.385	0.000	
0.142						
4	sub_grade	0.5339	0.007	77.851	0.000	
0.520						
5	emp_length	-0.0075	0.006	-1.265	0.206	-
0.019						
6	verification_status	-0.0087	0.006	-1.350	0.177	-
0.021						
7	dti	0.2370	0.006	39.629	0.000	
0.225						

8	initial_list_status	-0.0161	0.006	-2.707	0.007	-
0.028						
9	NONE	0.0057	0.005	1.079	0.281	-
0.005						
10	OWN	0.0578	0.006	9.434	0.000	
0.046						
11	RENT	0.1624	0.007	24.823	0.000	
0.150						
12	credit_card	0.0420	0.025	1.713	0.087	-
0.006						
13	debt_consolidation	0.0906	0.029	3.118	0.002	
0.034						
14	educational	0.0101	0.006	1.575	0.115	-
0.002						
15	home_improvement	0.0425	0.015	2.890	0.004	
0.014						
16	house	-0.0086	0.007	-1.154	0.249	-
0.023						
17	major_purchase	0.0165	0.011	1.513	0.130	-
0.005						
18	medical	0.0184	0.008	2.249	0.024	
0.002						
19	moving	0.0176	0.008	2.285	0.022	
0.003						
20	other	0.0248	0.014	1.740	0.082	-
0.003						
21	renewable_energy	0.0136	0.006	2.424	0.015	
0.003						
22	small_business	0.0688	0.009	7.514	0.000	
0.051						
23	vacation	0.0071	0.007	0.962	0.336	-
0.007						
24	wedding	-0.0391	0.009	-4.554	0.000	-
0.056						
25	JOINT	-0.0335	0.007	-4.780	0.000	-
0.047						

	upper_crit_val
1	-1.548
2	0.042
3	0.167
4	0.547
5	0.004
6	0.004
7	0.249
8	-0.004
9	0.016
10	0.070
11	0.175

12	0.090
13	0.148
14	0.023
15	0.071
16	0.006
17	0.038
18	0.034
19	0.033
20	0.053
21	0.025
22	0.087
23	0.022
24	-0.022
25	-0.020

```
logit_summary_2.columns=logit_summary.columns
```

```
logit_summary_2
```

	feature	coef	std err	z	pval	
lower_crit_val \						
1	const	-1.5604	0.006	-247.092	0.000	-
1.573						
2	loan_amnt	0.0279	0.007	4.031	0.000	
0.014						
3	term	0.1545	0.007	23.385	0.000	
0.142						
4	sub_grade	0.5339	0.007	77.851	0.000	
0.520						
5	emp_length	-0.0075	0.006	-1.265	0.206	-
0.019						
6	verification_status	-0.0087	0.006	-1.350	0.177	-
0.021						
7	dti	0.2370	0.006	39.629	0.000	
0.225						
8	initial_list_status	-0.0161	0.006	-2.707	0.007	-
0.028						
9	NONE	0.0057	0.005	1.079	0.281	-
0.005						
10	OWN	0.0578	0.006	9.434	0.000	
0.046						
11	RENT	0.1624	0.007	24.823	0.000	
0.150						
12	credit_card	0.0420	0.025	1.713	0.087	-
0.006						
13	debt_consolidation	0.0906	0.029	3.118	0.002	
0.034						
14	educational	0.0101	0.006	1.575	0.115	-
0.002						
15	home_improvement	0.0425	0.015	2.890	0.004	

0.014						
16	house	-0.0086	0.007	-1.154	0.249	-
0.023						
17	major_purchase	0.0165	0.011	1.513	0.130	-
0.005						
18	medical	0.0184	0.008	2.249	0.024	
0.002						
19	moving	0.0176	0.008	2.285	0.022	
0.003						
20	other	0.0248	0.014	1.740	0.082	-
0.003						
21	renewable_energy	0.0136	0.006	2.424	0.015	
0.003						
22	small_business	0.0688	0.009	7.514	0.000	
0.051						
23	vacation	0.0071	0.007	0.962	0.336	-
0.007						
24	wedding	-0.0391	0.009	-4.554	0.000	-
0.056						
25	JOINT	-0.0335	0.007	-4.780	0.000	-
0.047						

	upper_crit_val
1	-1.548
2	0.042
3	0.167
4	0.547
5	0.004
6	0.004
7	0.249
8	-0.004
9	0.016
10	0.070
11	0.175
12	0.090
13	0.148
14	0.023
15	0.071
16	0.006
17	0.038
18	0.034
19	0.033
20	0.053
21	0.025
22	0.087
23	0.022
24	-0.022
25	-0.020

logit_summary.columns

```
Index(['feature', 'coef', 'std err', 'z', 'pval', 'lower_crit_val',
      'upper_crit_val'],
      dtype='object')
```

```
logit_summary_2.columns=['feature', 'coef', 'std err', 'z', 'pval',
                        'lower_crit_val',
                        'upper_crit_val']
```

```
logit_summary_2[logit_summary_2.pval>0.05]
```

	feature	coef	std err	z	pval	lower_crit_val
5	emp_length	-0.0075	0.006	-1.265	0.206	-0.019
6	verification_status	-0.0087	0.006	-1.350	0.177	-0.021
9	NONE	0.0057	0.005	1.079	0.281	-0.005
12	credit_card	0.0420	0.025	1.713	0.087	-0.006
14	educational	0.0101	0.006	1.575	0.115	-0.002
16	house	-0.0086	0.007	-1.154	0.249	-0.023
17	major_purchase	0.0165	0.011	1.513	0.130	-0.005
20	other	0.0248	0.014	1.740	0.082	-0.003
23	vacation	0.0071	0.007	0.962	0.336	-0.007

	upper_crit_val
5	0.004
6	0.004
9	0.016
12	0.090
14	0.023
16	0.006
17	0.038
20	0.053
23	0.022

```
x_new_5=x_new_4.drop('vacation', axis=1)
```

```
xtrain_new_2, xtest_new_2, ytrain, ytest=train_test_split(x_new_5, y,
test_size=0.2, random_state=10)
```

```
xtrain_sc_new_2=sc.fit_transform(xtrain_new_2)
xtest_sc_new_2=sc.transform(xtest_new_2)
```

```
xtrain_sc_new_2=pd.DataFrame(xtrain_sc_new_2,
columns=xtrain_new_2.columns)
```

```
xtest_sc_new_2=pd.DataFrame(xtest_sc_new_2,
columns=xtest_new_2.columns)

xtrain_sc_new_2=sm.add_constant(xtrain_sc_new_2)

model_logit_3=sm.Logit(list(ytrain), xtrain_sc_new_2).fit()

Optimization terminated successfully.
      Current function value: 0.453418
      Iterations 6
```

```
model_logit_3.summary()
```

```
<class 'statsmodels.iolib.summary.Summary'>
"""
```

Logit Regression Results

```
=====
Dep. Variable:          y      No. Observations:
201350
Model:                Logit      Df Residuals:
201326
Method:                MLE      Df Model:
23
Date:                Fri, 18 Apr 2025      Pseudo R-squ.:
0.08554
Time:                19:10:50      Log-Likelihood:
-91296.
converged:                True      LL-Null:
-99836.
Covariance Type:                nonrobust      LLR p-value:
0.000
=====
```

		coef	std err	z	P> z	
[0.025 0.975]						

const		-1.5604	0.006	-247.100	0.000	-
1.573	-1.548					
loan_amnt		0.0278	0.007	4.011	0.000	
0.014	0.041					
term		0.1543	0.007	23.367	0.000	
0.141	0.167					
sub_grade		0.5342	0.007	77.968	0.000	
0.521	0.548					
emp_length		-0.0074	0.006	-1.248	0.212	-
0.019	0.004					
verification_status		-0.0086	0.006	-1.339	0.180	-

0.021	0.004					
dti		0.2371	0.006	39.645	0.000	
0.225	0.249					
initial_list_status		-0.0162	0.006	-2.718	0.007	-
0.028	-0.005					
NONE		0.0057	0.005	1.079	0.281	-
0.005	0.016					
OWN		0.0578	0.006	9.429	0.000	
0.046	0.070					
RENT		0.1624	0.007	24.822	0.000	
0.150	0.175					
credit_card		0.0285	0.020	1.432	0.152	-
0.011	0.067					
debt_consolidation		0.0742	0.023	3.187	0.001	
0.029	0.120					
educational		0.0091	0.006	1.443	0.149	-
0.003	0.022					
home_improvement		0.0349	0.012	2.837	0.005	
0.011	0.059					
house		-0.0113	0.007	-1.617	0.106	-
0.025	0.002					
major_purchase		0.0114	0.009	1.203	0.229	-
0.007	0.030					
medical		0.0150	0.007	2.042	0.041	
0.001	0.029					
moving		0.0145	0.007	2.081	0.037	
0.001	0.028					
other		0.0173	0.012	1.462	0.144	-
0.006	0.040					
renewable_energy		0.0126	0.006	2.287	0.022	
0.002	0.023					
small_business		0.0645	0.008	8.121	0.000	
0.049	0.080					
wedding		-0.0417	0.008	-5.121	0.000	-
0.058	-0.026					
JOINT		-0.0335	0.007	-4.780	0.000	-
0.047	-0.020					

=====

=====

"""

```
x_new_6=x_new_5.drop(['NONE','emp_length','major_purchase','verification_status', 'educational',
'credit_card','other','medical','moving','renewable_energy'], axis=1)
```

```
xtrain_new_3, xtest_new_3, ytrain, ytest=train_test_split(x_new_6, y,
test_size=0.2, random_state=10)
```

```
xtrain_sc_3=sc.fit_transform(xtrain_new_3)
xtest_sc_3=sc.transform(xtest_new_3)

xtrain_sc_3=pd.DataFrame(xtrain_sc_3, columns=xtrain_new_3.columns)
xtest_sc_3=pd.DataFrame(xtest_sc_3, columns=xtest_new_3.columns)

xtrain_sc_3=sm.add_constant(xtrain_sc_3)

model_logit_3=sm.Logit(list(ytrain), xtrain_sc_3).fit()

Optimization terminated successfully.
      Current function value: 0.453457
      Iterations 6
```

```
model_logit_3.summary()
```

```
<class 'statsmodels.iolib.summary.Summary'>
```

```
"""
```

Logit Regression Results

```
=====
=====
```

```
Dep. Variable:                y      No. Observations:
201350
Model:                Logit      Df Residuals:
201336
Method:                MLE      Df Model:
13
Date:                Fri, 18 Apr 2025      Pseudo R-squ.:
0.08546
Time:                19:10:51      Log-Likelihood:
-91304.
converged:                True      LL-Null:
-99836.
Covariance Type:                nonrobust      LLR p-value:
0.000
```

```
=====
=====
```

		coef	std err	z	P> z	
[0.025	0.975]					

const		-1.5602	0.006	-247.113	0.000	-
1.573	-1.548					
loan_amnt		0.0246	0.007	3.751	0.000	
0.012	0.037					
term		0.1533	0.007	23.265	0.000	
0.140	0.166					
sub_grade		0.5345	0.007	80.568	0.000	
0.522	0.548					
dti		0.2359	0.006	39.856	0.000	

0.224	0.248					
initial_list_status		-0.0166	0.006	-2.787	0.005	-
0.028	-0.005					
OWN		0.0579	0.006	9.466	0.000	
0.046	0.070					
RENT		0.1634	0.006	25.176	0.000	
0.151	0.176					
debt_consolidation		0.0372	0.007	5.715	0.000	
0.024	0.050					
home_improvement		0.0176	0.007	2.654	0.008	
0.005	0.031					
house		-0.0173	0.006	-2.875	0.004	-
0.029	-0.005					
small_business		0.0544	0.005	10.076	0.000	
0.044	0.065					
wedding		-0.0475	0.007	-6.453	0.000	-
0.062	-0.033					
JOINT		-0.0334	0.007	-4.775	0.000	-
0.047	-0.020					

```
=====
=====
"""
```

#Now that it is safe to use this, I will test

```
xtest_sc_3=sm.add_constant(xtest_sc_3)
```

```
ypred=model_logit_3.predict(xtest_sc_3)
```

```
from sklearn.metrics import classification_report, confusion_matrix,
roc_curve, roc_auc_score
```

```
ypred
```

```
0      0.288408
1      0.141869
2      0.193979
3      0.546137
4      0.201468
```

```
...
50333  0.153809
50334  0.229770
50335  0.130264
50336  0.194442
50337  0.101491
```

```
Length: 50338, dtype: float64
```

```
ypred_bin=[1 if i>=0.5 else 0 for i in ypred] #Putting a threshold of 0.5
```

```
print(confusion_matrix(ytest, ypred_bin))
```

```
[[39729 800]
 [ 9046 763]]
```

```
print(classification_report(ytest, ypred_bin)) #we see that it is
increasing the recall of my 0 class
```

	precision	recall	f1-score	support
0	0.81	0.98	0.89	40529
1	0.49	0.08	0.13	9809
accuracy			0.80	50338
macro avg	0.65	0.53	0.51	50338
weighted avg	0.75	0.80	0.74	50338

#however my aim is to predict those who were charged away-if I get them wrong the bank loses a lot of money

#therefore I want to increase the recall of my class 1 which is "charged away", I want to predict all of them correctly

```
ypred_bin=[1 if i>0.05 else 0 for i in ypred]
```

```
print(classification_report(ytest, ypred_bin))
#this model shows that while my recall of class 1 is now 100%,
#I am likely to lose a large number of good paying customers because I
misclassify them as people who are charged off
```

	precision	recall	f1-score	support
0	0.97	0.01	0.03	40529
1	0.20	1.00	0.33	9809
accuracy			0.21	50338
macro avg	0.58	0.51	0.18	50338
weighted avg	0.82	0.21	0.09	50338

```
fpr, tpr, thresh=roc_curve(ytest, ypred)
```

```
fpr
```

```
array([0.          , 0.          , 0.          , ..., 0.99962989,
0.99962989,
       1.          ])
```

```
tpr
```

```
array([0.00000000e+00, 1.01947191e-04, 2.03894383e-04, ...,
       9.99898053e-01, 1.00000000e+00, 1.00000000e+00])
```




```
pd.crosstab(ytest, ypred_bin)
```

```
col_0      0      1  
loan_status  
0         540  39989  
1          16   9793
```

```
#confusion_matrix(ytest, ypred_bin)
```

```
plt.plot(fpr,tpr)  
plt.plot([0,1],[0,1])  
plt.legend(['My model','Random model'])  
plt.xlabel('tpr')  
plt.ylabel('fpr')  
plt.title('ROC Curve shows that my model is better than a random  
model')  
plt.show()
```



```
roc_auc_score(ytest, ypred)
```

```
0.7027209058659821
```

```
y.value_counts() #This shows that the class which I want to predict is having very few rows
```

```
loan_status
0      202268
1       49420
Name: count, dtype: int64
```

```
#So we need to use SMOTE here to do oversampling
```

```
xtrain_sc_new_3=sc.fit_transform(xtrain_new_3)
```

```
xtest_sc_new_3=sc.transform(xtest_new_3)
```

```
import imblearn
```

```
from imblearn.over_sampling import SMOTE
```

```
smote=SMOTE()
```

```
xtrain_resample, ytrain_resample=smote.fit_resample(xtrain_sc_new_3, ytrain)
```

```
model_logit_sm=sm.Logit(ytrain_resample, xtrain_resample).fit()
```

```
Optimization terminated successfully.
      Current function value: 0.629787
      Iterations 5
```

```
ypred_smote=model_logit_sm.predict(xtest_sc_new_3)
```

```
ypred_smote_bin=[1 if i>=0.5 else 0 for i in ypred_smote]
```

```
print(classification_report(ytest, ypred_smote_bin))
```

	precision	recall	f1-score	support
0	0.89	0.60	0.72	40529
1	0.30	0.69	0.41	9809
accuracy			0.62	50338
macro avg	0.59	0.65	0.57	50338
weighted avg	0.77	0.62	0.66	50338

```
pd.crosstab(ytest, ypred_smote_bin) #this shows our recall for the class 1 has gone up from 8% to 69% by ensuring that our classes are balanced
```

col_0	0	1
loan_status		
0	24342	16187
1	3014	6795

```
#To increase the recall further I am reducing the threshold
ypred_smote_bin=[1 if i>=0.2 else 0 for i in ypred_smote]
```

```
print(classification_report(ytest, ypred_smote_bin))
```

	precision	recall	f1-score	support
0	0.97	0.01	0.02	40529
1	0.20	1.00	0.33	9809
accuracy			0.20	50338
macro avg	0.58	0.51	0.18	50338
weighted avg	0.82	0.20	0.08	50338

#comparing the scores for the predictions after smote (above) what is below, we see that the recall has increased for the majority class with the same threshold


```
ypred_bin_copy=[1 if i>=0.5 else 0 for i in ypred] #replicating the older model predictions with 0.5 threshold just for comparison
```

```
print(classification_report(ytest, ypred_bin_copy))
```

	precision	recall	f1-score	support
0	0.81	0.98	0.89	40529
1	0.49	0.08	0.13	9809
accuracy			0.80	50338
macro avg	0.65	0.53	0.51	50338
weighted avg	0.75	0.80	0.74	50338

```
fpr_sm, tpr_sm, thresh=roc_curve(ytest, ypred_smote)
fpr, tpr, thresh=roc_curve(ytest, ypred)
```

```
plt.plot(fpr_sm, tpr_sm)
plt.plot(fpr, tpr)
plt.legend(['After SMOTE', 'Before oversampling'])
plt.title('The score is almost identical before and after oversampling')
plt.xlabel('fpr')
plt.ylabel('tpr')
plt.show()
```



```
roc_auc_score(ytest, ypred_smote) #this is after smote
0.7029267220245609

roc_auc_score(ytest, ypred) #this is before oversampling
0.7027209058659821

#therefore, I see a very small increase in the overall performance of
the model even after oversampling and balancing
#At a threshold of 0.5 however, the increase in recall is nicely
visible because of oversampling.

#trying to see if the difference can better seen using the precision
recall score


from sklearn.metrics import precision_recall_curve

pr, re, thr=precision_recall_curve(ytest, ypred) #This is for the
model without oversampling
pr_sm, re_sm, thr=precision_recall_curve(ytest, ypred_smote) #this is
after smote

plt.plot(re, pr, label='logistic')
plt.plot(re_sm, pr_sm, label='logistic-sm')
```

```
plt.legend(['Before oversampling', 'After oversampling'])
plt.xlabel('recall')
plt.ylabel('precision')
plt.title('Precision recall curve shows no great difference between
the unbalanced and oversampled model')
plt.show()
```

Precision recall curve shows no great difference between the unbalanced and oversampled model



```
from sklearn.metrics import auc
```

```
print(auc(re, pr))
```

```
0.35505049832411795
```

```
print(auc(re_sm, pr_sm))
```

```
0.35423690129110513
```

#Both these show that my model is actually worse than an average model.

#This shows something wrong with my data itself since my precision recall score should be above 0.5 which is for a random model.

#Conclusions and recommendations

#1. More loans are going to b and c grade people when compared to a grade people and this increases the risk

#Therefore, there is an opportunity for the company to reroute most of its loan amount to customers with A grade

#2. #People who paid up had higher incomes than people who did not,

and our data shows that most of the loans go to people with very low incomes.

#So the company has an opportunity to identify factors other than income which help predict whether a person pays up, and then use that to qualify low income groups for a loan

#3.Data that on an average a higher loan is being provided to those in the G category who have the highest risk of default.

#This means the company should lower the risk of losing money, instead providing the highest loans to grade A customers and smaller loan to riskier customers

#4. The bank gives more loans to riskier categories with high interest: particularly restructuring previous loans or credit cards

#If we want to reduce the risk of default, the company should give smaller loans at lower interest to people in safer categories,

#such as for education or for household maintenance

#So the company should aim to give loans in safer categories, instead of giving more loans to a few high risk categories.

#Since weddings, cars, and education are the categories with lowest default rates,

#the company should give smaller loans to more people with lower interest rates in these categories

#5 #the company gives high interest rates to risky categories, and penalizes innocent borrowers to account for those who are defaulting

#The most likely reason is because they do not maintain data at an individual level and adjust it regularly based on each customer's behaviour

#Therefore in this case, I recommend that the bank should implement a blockchain solution

#so that they can identify who is likely to default rather than just which category is at high risk using aggregate data

#The most common way to do this is to trace the purchase behaviour and patterns

#from say the retail stores or supermarkets or online shopping for the same customers by partnering with retail or credit card companies to analyse behaviour across

#6. The overall number of people taking a loan reduces as their experience level increases, however the number of people defaulting remains the same

#Therefore the proportion of defaulting people increases in case of more experienced people.

#Therefore I recommend increasing interest rates for people who higher level of experience

#because the average loan ticket is the same as for other groups, yet default rate is more because fewer people take a loan overall

After feature selection and checking for multicollinearity, the most important features to predict the loan status accurately are:

a) sub grade

b) DTI

c) RENT (Whether the person is taking the loan for paying rent)

d) Term of the loan in months

#This dataset was unbalanced so using oversampling using SMOTE helped me improve the recall of my class 1 (the defaulter) from 8% to 69% at a threshold of 0.5

however it did not improve the overall performance of the model from the roc_auc_score

#Since my objective is to maximize the recall of my class 1 (people who defaulted), I reduced my threshold and it became 100%.

#However in this case I am then getting a very high degree of false positive errors with a recall of class 0 of only 1% and we can end up losing genuinely paying customers

#Also as the threshold increases my SMOTE does not have any impact on the prediction in this dataset,

#the impact is only visible when I keep my threshold at 0.5

#Also for unbalanced data like this, I read that precision recall curves are better and provide more information.

#In this case the score shows below 0.5 indicating that my model performs worse than an average model would.

#However my ROC curve shows a 70% performance.

#Therefore, my understanding is that I have not done enough feature engineering or identified the right feature combination, because very few of the important features had a good p value of <0.05 in the summary.

#Therefore, I will request the business to give me other behavioural data or purchase data from other sources, or maybe data on assets,

#that has not been modified by the company, yet which influences the outcome, because such variables will not have multicollinearity and be independent.