

```
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
```

```
import datetime as dt
```

```
data=pd.read_csv('ola_driver_scaler.csv')
```

```
data.head()
```

	Unnamed: 0	MMM-YY	Driver_ID	Age	Gender	City	Education_Level
0	0	01/01/19	1	28.0	0.0	C23	2
1	1	02/01/19	1	28.0	0.0	C23	2
2	2	03/01/19	1	28.0	0.0	C23	2
3	3	11/01/20	2	31.0	0.0	C7	2
4	4	12/01/20	2	31.0	0.0	C7	2

	Income	Dateofjoining	LastWorkingDate	Joining	Designation	Grade
0	57387	24/12/18	NaN			1
1	57387	24/12/18	NaN			1
2	57387	24/12/18	03/11/19			1
3	67016	11/06/20	NaN			2
4	67016	11/06/20	NaN			2

	Total Business Value	Quarterly Rating
0	2381060	2
1	-665480	2
2	0	2
3	0	1
4	0	1

```
data.shape
```

```
(19104, 14)
```

```
data.drop('Unnamed: 0', axis=1, inplace=True)
```

```
data.columns
```

```
Index(['MMM-YY', 'Driver_ID', 'Age', 'Gender', 'City',
      'Education_Level',
      'Income', 'Dateofjoining', 'LastWorkingDate', 'Joining
Designation',
      'Grade', 'Total Business Value', 'Quarterly Rating'],
      dtype='object')
```

```
data['Attrition']=[0 if pd.isna(i) else 1 for i in  
data['LastWorkingDate']] #Create the new target column
```

```
data['Attrition'].value_counts() #This data is unbalanced
```

```
Attrition  
0      17488  
1       1616  
Name: count, dtype: int64
```

#Given that we know it is important to retain drivers rather than hire new ones, it is important that we predict the class 1 correctly

```
data.dtypes
```

MMM-YY	object
Driver_ID	int64
Age	float64
Gender	float64
City	object
Education_Level	int64
Income	int64
Dateofjoining	object
LastWorkingDate	object
Joining Designation	int64
Grade	int64
Total Business Value	int64
Quarterly Rating	int64
Attrition	int64
dtype:	object

```
data.isnull().sum()
```

MMM-YY	0
Driver_ID	0
Age	61
Gender	52
City	0
Education_Level	0
Income	0
Dateofjoining	0
LastWorkingDate	17488
Joining Designation	0
Grade	0
Total Business Value	0
Quarterly Rating	0
Attrition	0
dtype:	int64

```
data['DOJ']=pd.to_datetime(data['Dateofjoining'], format='mixed')
```

```
data['LWD']=pd.to_datetime(data['LastWorkingDate'], format='mixed')
```

```
data.shape
```

```
(19104, 16)
```

```
data.head()
```

	MMM-YY	Driver_ID	Age	Gender	City	Education_Level	Income	\
0	01/01/19	1	28.0	0.0	C23	2	57387	
1	02/01/19	1	28.0	0.0	C23	2	57387	
2	03/01/19	1	28.0	0.0	C23	2	57387	
3	11/01/20	2	31.0	0.0	C7	2	67016	
4	12/01/20	2	31.0	0.0	C7	2	67016	

	Dateofjoining	LastWorkingDate	Joining	Designation	Grade	\
0	24/12/18	NaN			1	1
1	24/12/18	NaN			1	1
2	24/12/18	03/11/19			1	1
3	11/06/20	NaN			2	2
4	11/06/20	NaN			2	2

	Total Business Value	Quarterly Rating	Attrition	DOJ
LWD				
0	2381060	2	0	2018-12-24
NaT				
1	-665480	2	0	2018-12-24
NaT				
2	0	2	1	2018-12-24 2019-03-11
3	0	1	0	2020-11-06
NaT				
4	0	1	0	2020-11-06
NaT				

```
data['LWD'].head()
```

0	NaT
1	NaT
2	2019-03-11
3	NaT
4	NaT

```
Name: LWD, dtype: datetime64[ns]
```

```
data.drop(['Dateofjoining', 'LastWorkingDate'], axis=1, inplace=True)
```

```
data.rename(columns={'DOJ': 'Dateofjoining', 'LWD': 'LastWorkingDate'},  
inplace=True)
```

```
data.groupby('Joining Designation')
['Attrition'].sum()/data.groupby('Joining Designation')
['Attrition'].count()
```

```
Joining Designation
1    0.076493
2    0.094039
3    0.096242
4    0.064516
5    0.061538
Name: Attrition, dtype: float64
```

```
tenure_days=[]
for i in range(data.shape[0]):
    if (data.loc[i,'LastWorkingDate']) is pd.NaT:
        tenure_days.append(pd.to_datetime(dt.date.today())-
data.loc[i,'Dateofjoining'])
    else:
        tenure_days.append(data.loc[i,'LastWorkingDate']-
data.loc[i,'Dateofjoining'])
```

```
tenure_days
```

```
[Timedelta('2341 days 00:00:00'),
Timedelta('2341 days 00:00:00'),
Timedelta('77 days 00:00:00'),
Timedelta('1658 days 00:00:00'),
Timedelta('1658 days 00:00:00'),
Timedelta('1993 days 00:00:00'),
Timedelta('1993 days 00:00:00'),
Timedelta('1993 days 00:00:00'),
Timedelta('1993 days 00:00:00'),
Timedelta('142 days 00:00:00'),
Timedelta('2325 days 00:00:00'),
Timedelta('2325 days 00:00:00'),
Timedelta('57 days 00:00:00'),
Timedelta('1756 days 00:00:00'),
Timedelta('1756 days 00:00:00'),
Timedelta('1756 days 00:00:00'),
Timedelta('1756 days 00:00:00'),
Timedelta('1756 days 00:00:00'),
Timedelta('1706 days 00:00:00'),
Timedelta('1706 days 00:00:00'),
Timedelta('57 days 00:00:00'),
Timedelta('1627 days 00:00:00'),
Timedelta('2154 days 00:00:00'),
Timedelta('2154 days 00:00:00'),
Timedelta('2154 days 00:00:00')]
```

[illegible]

[illegible]

[illegible]

[illegible]

[illegible]

[illegible]

[illegible]

[illegible]

[illegible]

[illegible]

[illegible]

[illegible]

[illegible]

[illegible]

[illegible]

[illegible]

[illegible]

```
Timedelta('2014 days 00:00:00'),
Timedelta('2014 days 00:00:00'),
Timedelta('2014 days 00:00:00'),
Timedelta('2014 days 00:00:00'),
Timedelta('129 days 00:00:00'),
Timedelta('2063 days 00:00:00'),
Timedelta('2063 days 00:00:00'),
Timedelta('2063 days 00:00:00'),
Timedelta('2063 days 00:00:00'),
Timedelta('2063 days 00:00:00'),
Timedelta('2063 days 00:00:00'),
Timedelta('183 days 00:00:00'),
Timedelta('1902 days 00:00:00'),
Timedelta('1902 days 00:00:00'),
Timedelta('78 days 00:00:00'),
Timedelta('1749 days 00:00:00'),
Timedelta('1749 days 00:00:00'),
Timedelta('1749 days 00:00:00'),
Timedelta('1749 days 00:00:00'),
Timedelta('1749 days 00:00:00'),
Timedelta('2419 days 00:00:00'),
Timedelta('2419 days 00:00:00'),
Timedelta('2419 days 00:00:00'),
Timedelta('2419 days 00:00:00'),
Timedelta('2419 days 00:00:00'),
Timedelta('2419 days 00:00:00'),
Timedelta('2419 days 00:00:00'),
Timedelta('351 days 00:00:00'),
Timedelta('2489 days 00:00:00'),
Timedelta('2489 days 00:00:00'),
Timedelta('2489 days 00:00:00'),
Timedelta('2489 days 00:00:00'),
Timedelta('2489 days 00:00:00'),
Timedelta('2489 days 00:00:00'),
Timedelta('357 days 00:00:00'),
Timedelta('2575 days 00:00:00'),
Timedelta('2575 days 00:00:00'),
Timedelta('301 days 00:00:00'),
Timedelta('1776 days 00:00:00'),
Timedelta('1776 days 00:00:00'),
Timedelta('1776 days 00:00:00'),
Timedelta('1776 days 00:00:00'),
Timedelta('1776 days 00:00:00'),
Timedelta('1952 days 00:00:00'),
Timedelta('1952 days 00:00:00'),
Timedelta('1952 days 00:00:00'),
Timedelta('1952 days 00:00:00'),
```

[illegible]

```
Timedelta('3310 days 00:00:00'),
Timedelta('3310 days 00:00:00'),
Timedelta('3310 days 00:00:00'),
Timedelta('3310 days 00:00:00'),
Timedelta('3310 days 00:00:00'),
Timedelta('3310 days 00:00:00'),
Timedelta('3310 days 00:00:00'),
Timedelta('3310 days 00:00:00'),
Timedelta('3310 days 00:00:00'),
Timedelta('3310 days 00:00:00'),
Timedelta('3310 days 00:00:00'),
Timedelta('3310 days 00:00:00'),
Timedelta('3310 days 00:00:00'),
Timedelta('3310 days 00:00:00'),
Timedelta('3310 days 00:00:00'),
Timedelta('3310 days 00:00:00'),
Timedelta('1777 days 00:00:00'),
Timedelta('1777 days 00:00:00'),
Timedelta('1777 days 00:00:00'),
Timedelta('1777 days 00:00:00'),
Timedelta('113 days 00:00:00'),
Timedelta('2538 days 00:00:00'),
Timedelta('243 days 00:00:00'),
Timedelta('2578 days 00:00:00'),
Timedelta('2578 days 00:00:00'),
Timedelta('302 days 00:00:00'),
Timedelta('2103 days 00:00:00'),
Timedelta('2103 days 00:00:00'),
Timedelta('2103 days 00:00:00'),
Timedelta('97 days 00:00:00'),
Timedelta('1925 days 00:00:00'),
Timedelta('1925 days 00:00:00'),
Timedelta('1925 days 00:00:00'),
Timedelta('1925 days 00:00:00'),
Timedelta('130 days 00:00:00'),
Timedelta('2133 days 00:00:00'),
Timedelta('2133 days 00:00:00'),
Timedelta('2133 days 00:00:00'),
Timedelta('73 days 00:00:00'),
Timedelta('1623 days 00:00:00'),
Timedelta('1861 days 00:00:00'),
Timedelta('1861 days 00:00:00'),
Timedelta('1861 days 00:00:00'),
...]
```

```
data['tenure in days']=tenure_days
```

```
data['tenure in days']
```



```

0      2341 days
1      2341 days
2         77 days
3      1658 days
4      1658 days
...
19099   1809 days
19100   1809 days
19101   1809 days
19102   1809 days
19103   1809 days
Name: tenure in days, Length: 19104, dtype: timedelta64[ns]

#remove this if the cell above works well
#data['tenure in days']=data['LastWorkingDate']-data['Dateofjoining']

data.groupby('Attrition')['tenure in days'].mean() #the average tenure
of those who have left is almost 1 year

Attrition
0    2602 days 17:12:14.492223232
1     357 days 13:46:02.376237624
Name: tenure in days, dtype: timedelta64[ns]

data[data['Attrition']==1]['tenure in days']

2         77 days
9        142 days
12         57 days
20         57 days
27        175 days
...
19039     19 days
19054     92 days
19081     61 days
19090    418 days
19096    334 days
Name: tenure in days, Length: 1616, dtype: timedelta64[ns]

data.rename(columns={'MMM-YY':'Reporting_date'}, inplace=True)

data.dtypes

Reporting_date      object
Driver_ID           int64
Age                float64
Gender             float64
City               object
Education_Level     int64

```

```

Income                int64
Joining Designation   int64
Grade                 int64
Total Business Value  int64
Quarterly Rating      int64
Attrition             int64
Dateofjoining         datetime64[ns]
LastWorkingDate       datetime64[ns]
tenure in days        timedelta64[ns]
dtype: object

```

```

data['Reporting_date']=pd.to_datetime(data['Reporting_date'],
format='mixed')

```

```

data.dtypes

```

```

Reporting_date        datetime64[ns]
Driver_ID             int64
Age                   float64
Gender                float64
City                  object
Education_Level        int64
Income                int64
Joining Designation    int64
Grade                 int64
Total Business Value  int64
Quarterly Rating      int64
Attrition             int64
Dateofjoining         datetime64[ns]
LastWorkingDate       datetime64[ns]
tenure in days        timedelta64[ns]
dtype: object

```

#performing knn imputation of on the age and gender columns

```

import sklearn
from sklearn.impute import KNNImputer, SimpleImputer

knn_imputer=KNNImputer(n_neighbors=5)
sim_imputer=SimpleImputer(strategy='most_frequent')

age=knn_imputer.fit_transform(pd.DataFrame(data['Age'])).reshape(1,-1)
gender=sim_imputer.fit_transform(pd.DataFrame(data['Gender'])).reshape(1,-1)

age=pd.Series(age[0])
gender=pd.Series(gender[0])

data['Age']=age

```

```

data['Gender']=gender
data.isnull().sum()

Reporting_date      0
Driver_ID           0
Age                0
Gender              0
City               0
Education_Level     0
Income              0
Joining Designation 0
Grade              0
Total Business Value 0
Quarterly Rating    0
Attrition           0
Dateofjoining       0
LastWorkingDate     17488
tenure in days      0
dtype: int64

data.groupby('Gender')['Attrition'].sum()/data.groupby('Gender')
['Attrition'].count() #Both male and female drivers have equal
attrition rate

Gender
0.0    0.085296
1.0    0.083605
Name: Attrition, dtype: float64


data.groupby(['Gender','Attrition'])['Total Business Value'].mean()

Gender  Attrition
0.0     0          611713.833153
        1          36629.789252
1.0     0          633163.740938
        1          47674.392804
Name: Total Business Value, dtype: float64

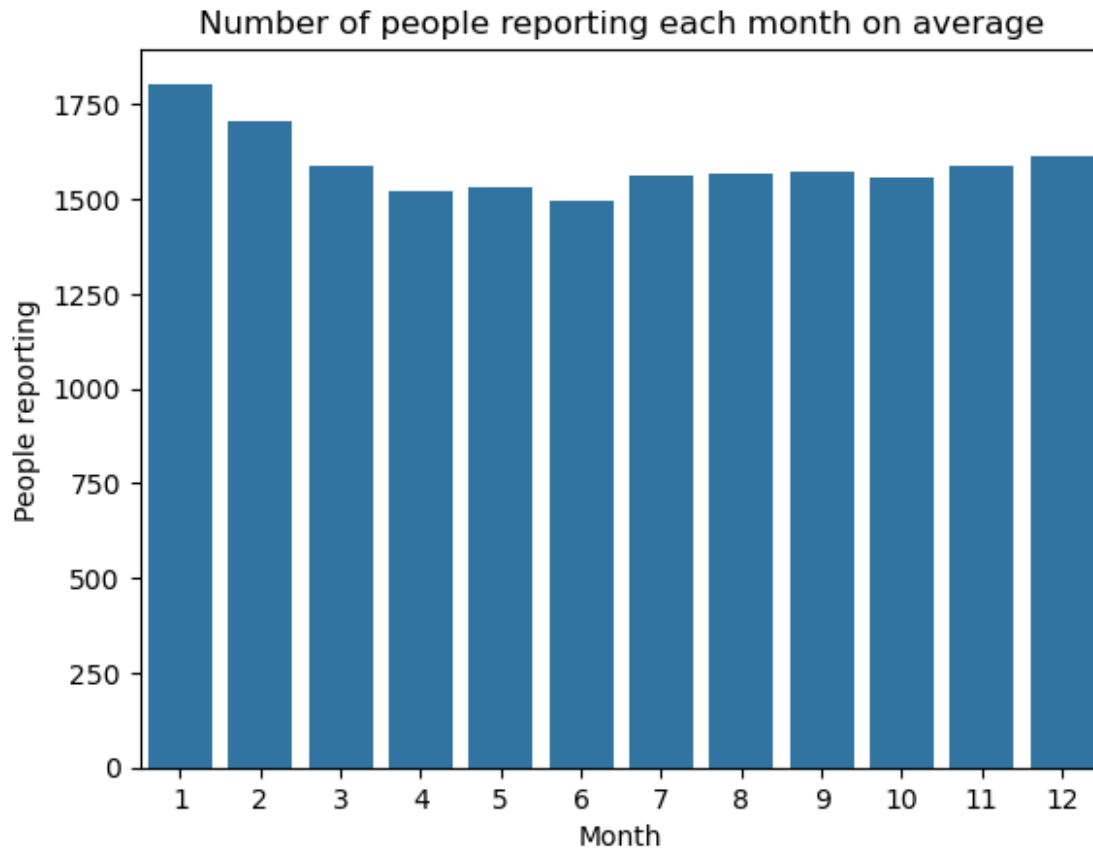
sns.barplot(data=data, x='Attrition', y='Total Business Value',
hue='Gender') #We see that women add more business value than men
whether they stay or leave

<Axes: xlabel='Attrition', ylabel='Total Business Value'>


```



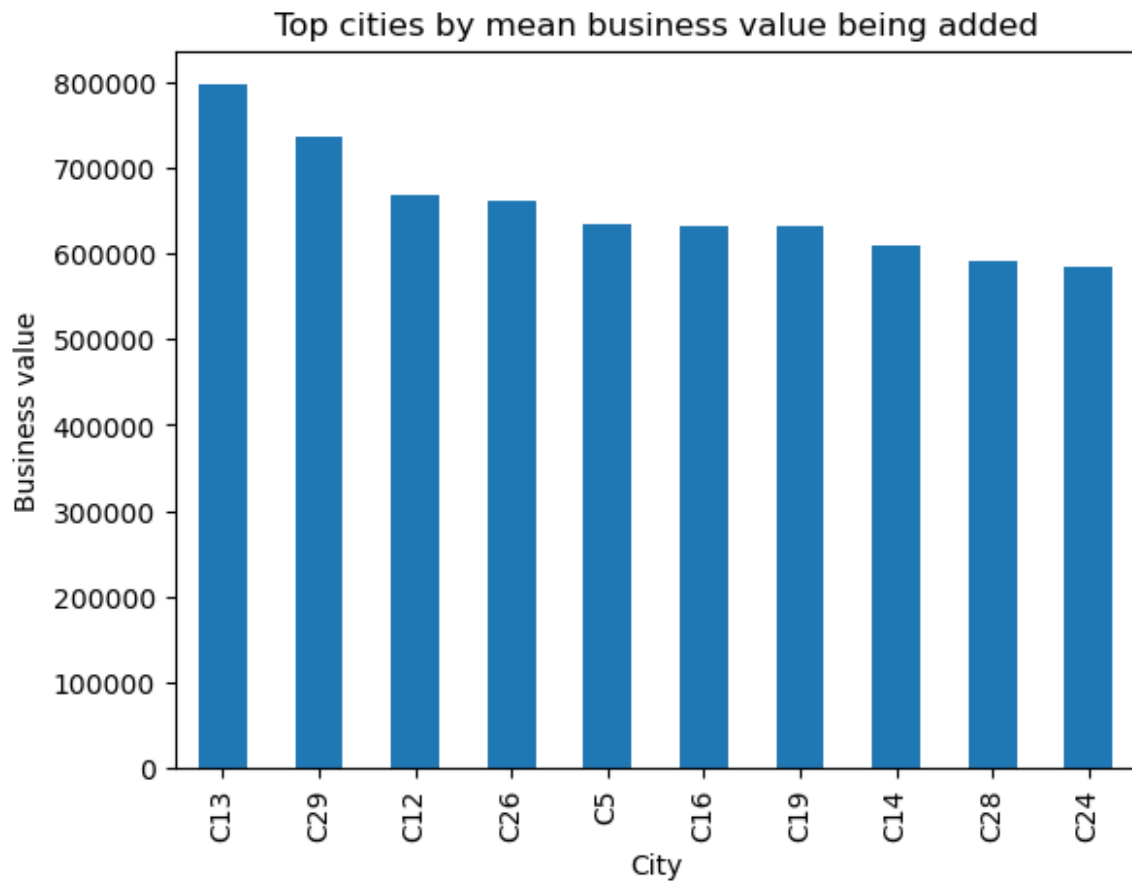
```
sns.countplot(x=data['Reporting_date'].dt.month)
plt.title('Number of people reporting each month on average')
plt.xlabel('Month')
plt.ylabel('People reporting')
plt.show()
```




```
data.groupby('City')['Income'].mean().sort_values(ascending=False)
[:10].plot(kind='bar')
plt.title('Top 10 cities by driver income')
plt.ylabel('Income')
plt.show()
```



```
data.groupby('City')['Total Business  
Value'].mean().sort_values(ascending=False)[:10].plot(kind='bar')  
plt.ylabel('Business value')  
plt.title('Top cities by mean business value being added')  
plt.show()
```




```
data.groupby('City')['Driver_ID'].count().sort_values(ascending=False)
[:10].plot(kind='bar')
plt.title('Top city by driver count')
Text(0.5, 1.0, 'Top city by driver count')
```




*#We see that the most attractive cities appearing in each of the 3 graphs are C12, C26, and C29
#So the company should look at putting demand and number of rides to decide how many drivers to hire in a each city.*

```
sns.barplot(data=data, x='Education_Level', y='Income', hue='Gender')  
<Axes: xlabel='Education_Level', ylabel='Income'>
```

```
sns.barplot(data=data, x='Education_Level',  
y='Income', hue='Attrition') #People who left have lower incomes than  
people who stayed
```

```
<Axes: xlabel='Education_Level', ylabel='Income'>
```



#More educated people get higher ratings in general-and in highest education, the men get more income on average than women

```
data.columns
```


```
Index(['Reporting_date', 'Driver_ID', 'Age', 'Gender', 'City',
      'Education_Level', 'Income', 'Joining Designation', 'Grade',
      'Total Business Value', 'Quarterly Rating', 'Attrition',
      'Dateofjoining', 'LastWorkingDate', 'tenure in days'],
      dtype='object')
```

```
data.groupby('Gender')['Attrition'].value_counts()
```

```
Gender  Attrition
0.0     0          10177
        1           949
1.0     0          7311
        1           667
Name: count, dtype: int64
```


```
data['Age'].hist()
```

```
<Axes: >
```




```
data['Income'].hist() #This is a little skewed so I can apply a log transformation, I will leave it as is for now
```

```
<Axes: >
```



```
sns.scatterplot(data=data, x='Age', y='Income', hue='Grade',  
palette='Set1') #We see that age is not correlated with grade, however  
it appears income is
```

```
<Axes: xlabel='Age', ylabel='Income'>
```



```
data['Attrition'].value_counts()
```

```
Attrition
```

```
0    17488
```


```
1     1616
```

```
Name: count, dtype: int64
```

```
data.groupby('Attrition')['Age'].plot(kind='kde')
```


```
plt.legend(['No attrition', 'Attrition'])
```

```
<matplotlib.legend.Legend at 0x165969768a0>
```



```
sns.countplot(hue='Attrition', x='Quarterly Rating', data=data,  
palette='Set1')  
plt.title('More people with quarterly ratings 1 and 2 are leaving  
compared to 3 and 4')  
plt.show()
```


More people with quarterly ratings 1 and 2 are leaving compared to 3 and 4



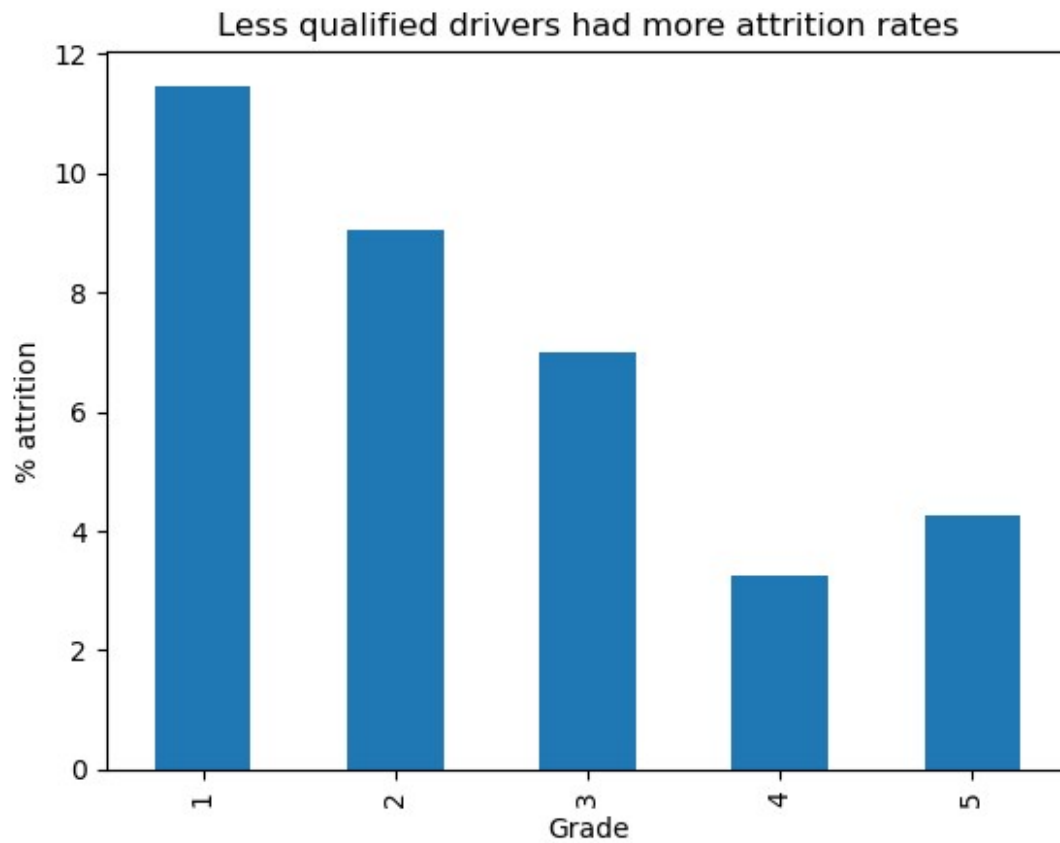
```
sum(data.groupby('Driver_ID')
     ['Reporting_date'].count().sort_values(ascending=False)==24) #229 of
the drivers reported the maximum duration of 24 days
```

229

```
data.groupby('Attrition')['Income'].plot(kind='kde')
plt.legend(['No attrition', 'Attrition'])
plt.show()
#People who left already had a higher income on average
```




```
(data.groupby('Grade')['Attrition'].sum()*100/data.groupby('Grade')
['Attrition'].count()).plot(kind='bar')
plt.ylabel('% attrition')
plt.title('Less qualified drivers had more attrition rates')
plt.show()
#This is more meaningful than the graph above and shows that attrition
is around 11% for grade 1 and only 4% of grade 5 drivers
```

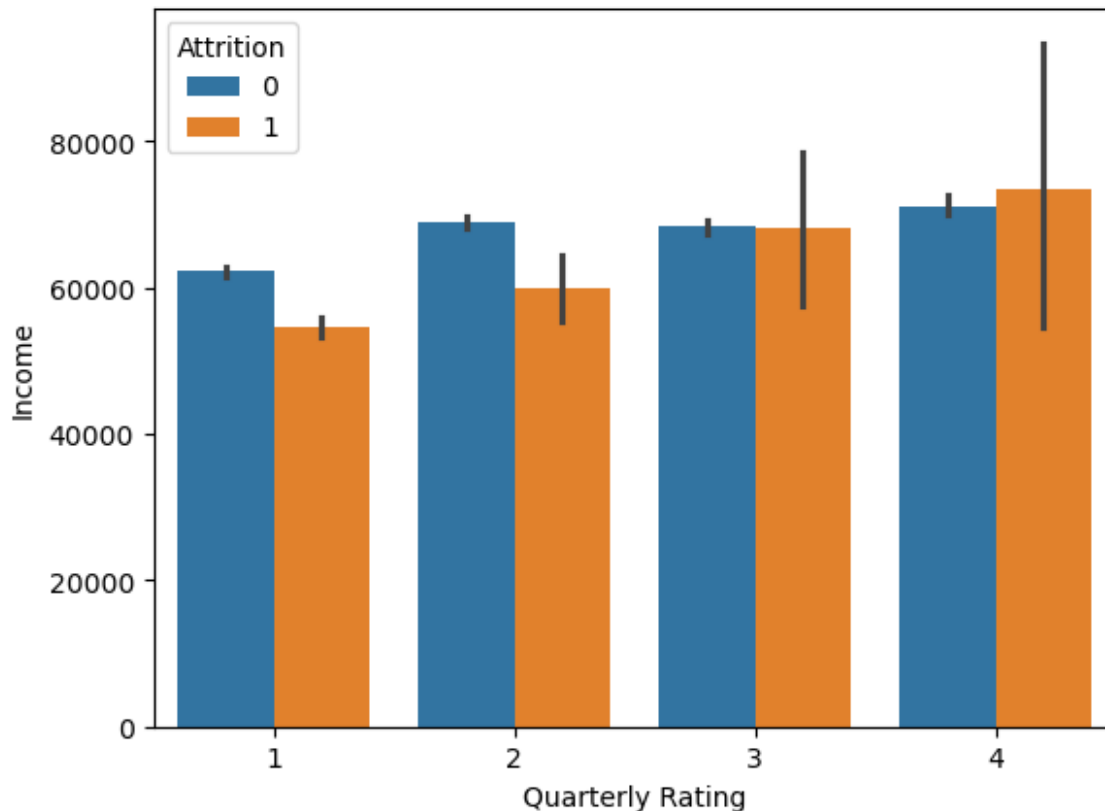
```
sns.barplot(x='Grade', y='Income', hue='Attrition', data=data)
plt.title('People who left already had a lower income compared to
those who did not')
plt.show()
```

People who left already had a lower income compared to those who did not



```
sns.barplot(data=data, x='Quarterly Rating', y='Income',  
hue='Attrition')
```

```
<Axes: xlabel='Quarterly Rating', ylabel='Income'>
```



#1. The quarterly rating does not impact the income of the drivers very significantly, at least among those leaving
#2. For those staying, we see larger error bars, which means the income cannot be predicted confidently as the quarterly rating increases
#3. At higher quarterly ratings of 3 or 4, the people who left and the people who stayed had the same mean income, the only difference is in the error bars and the confidence of the prediction
#For people in low quarterly ratings, the income of those who left was lower than those who stayed.
#However, for people with higher quarterly ratings people with similar or higher income are also leaving

```
data['tenure in days']
```

```
0    2341 days
1    2341 days
2      77 days
3   1658 days
```

```

4          1658 days
...
19099      1809 days
19100      1809 days
19101      1809 days
19102      1809 days
19103      1809 days
Name: tenure in days, Length: 19104, dtype: timedelta64[ns]

data['tenure in days']=data['tenure in days'].apply(str).str.split('
', expand=True)[0]


data['tenure in days']

0          2341
1          2341
2           77
3          1658
4          1658
...
19099      1809
19100      1809
19101      1809
19102      1809
19103      1809
Name: tenure in days, Length: 19104, dtype: object

sns.scatterplot(data=data, y='Income', x='Total Business Value',
hue='Attrition')
#Almost all the people who are leaving did not add any business value

<Axes: xlabel='Total Business Value', ylabel='Income'>

```



```
data.groupby('Attrition')['Total Business Value'].mean() #Those who left were adding lesser business value
```

```
Attrition
0      620681.140782
1       41188.422030
Name: Total Business Value, dtype: float64
```

```
#data.groupby('Driver_ID')['Total Business Value'].sum()<=0
```

```
#data['Driver_ID'].count()
```

```
 #(data['Total Business Value']>0).sum()
```

```
#data.groupby('Driver_ID')(['Total Business Value']<=0).sum()/data.groupby('Driver_ID')['Driver_ID'].count()
```

```
data.dtypes
```

```
Reporting_date      datetime64[ns]
Driver_ID           int64
Age                 float64
```

```
Gender float64
City object
Education_Level int64
Income int64
Joining Designation int64
Grade int64
Total Business Value int64
Quarterly Rating int64
Attrition int64
Dateofjoining datetime64[ns]
LastWorkingDate datetime64[ns]
tenure in days object
dtype: object
```

```
data_copy=data.drop(['Reporting_date',  
'Driver_ID','Dateofjoining','LastWorkingDate','tenure in days'],  
axis=1)
```

#I am dropping the tenure as well since it is causing overfitting

```
data_copy.dtypes
```

```
Age float64
Gender float64
City object
Education_Level int64
Income int64
Joining Designation int64
Grade int64
Total Business Value int64
Quarterly Rating int64
Attrition int64
dtype: object
```

```
data copy.City.unique()
```

```
array(['C23', 'C7', 'C13', 'C9', 'C11', 'C2', 'C19', 'C26', 'C20',  
      'C17', 'C29', 'C10', 'C24', 'C14', 'C6', 'C28', 'C5', 'C18', 'C27',  
      'C15', 'C8', 'C25', 'C21', 'C1', 'C4', 'C3', 'C16', 'C22', 'C12'],  
      dtype=object)
```

```
city=pd.get_dummies(data_copy['City'], drop_first=True).astype(int)
```

city

[illegible]

1	0	0	0	0	0	0	0	0	0	0	...	0	0
0													
2	0	0	0	0	0	0	0	0	0	0	...	0	0
0													
3	0	0	0	0	0	0	0	0	0	0	...	0	0
0													
4	0	0	0	0	0	0	0	0	0	0	...	0	0
0													
...
...													
19099	0	0	0	0	0	0	0	0	0	0	...	1	0
0													
19100	0	0	0	0	0	0	0	0	0	0	...	1	0
0													
19101	0	0	0	0	0	0	0	0	0	0	...	1	0
0													
19102	0	0	0	0	0	0	0	0	0	0	...	1	0
0													
19103	0	0	0	0	0	0	0	0	0	0	...	1	0
0													

	C3	C4	C5	C6	C7	C8	C9
0	0	0	0	0	0	0	0
1	0	0	0	0	0	0	0
2	0	0	0	0	0	0	0
3	0	0	0	0	1	0	0
4	0	0	0	0	1	0	0
...
19099	0	0	0	0	0	0	0
19100	0	0	0	0	0	0	0
19101	0	0	0	0	0	0	0
19102	0	0	0	0	0	0	0
19103	0	0	0	0	0	0	0

[19104 rows x 28 columns]

```
data_copy=pd.concat([data_copy.drop('City', axis=1), city],
axis=1).dropna()
```

```
import statsmodels
import statsmodels.api as sm

y=data_copy['Attrition']
x=data_copy.drop('Attrition', axis=1)
x.isnull().sum()
```

Age	0
Gender	0
Education_Level	0
Income	0
Joining Designation	0
Grade	0
Total Business Value	0
Quarterly Rating	0
C10	0
C11	0
C12	0
C13	0
C14	0
C15	0
C16	0
C17	0
C18	0
C19	0
C2	0
C20	0
C21	0
C22	0
C23	0
C24	0
C25	0
C26	0
C27	0
C28	0
C29	0
C3	0
C4	0
C5	0
C6	0
C7	0
C8	0
C9	0

dtype: int64

```
import sklearn
from sklearn.model_selection import train_test_split

xtrain, xtest, ytrain, ytest=train_test_split(x, y, test_size=0.2,
random_state=10)

#xtrain['tenure in days']=xtrain['tenure in days'].astype('int')

from sklearn.preprocessing import StandardScaler

sc=StandardScaler()
```



```

xtrain_num=xtrain[['Age','Income','Total Business Value']]
xtest_num=xtest[['Age','Income','Total Business Value']]

xtrain_sc_num=pd.DataFrame(sc.fit_transform(xtrain_num),
columns=['Age','Income','Total Business Value'])
xtest_sc_num=pd.DataFrame(sc.transform(xtest_num),
columns=['Age','Income','Total Business Value'])

xtrain_sc_num.index=xtrain.index
xtest_sc_num.index=xtest.index


xtrain_cat=xtrain.drop(['Age','Income','Total Business Value'],
axis=1)
xtest_cat=xtest.drop(['Age','Income','Total Business Value'], axis=1)


xtrain=pd.concat([xtrain_cat, xtrain_sc_num], axis=1)
xtest=pd.concat([xtest_cat, xtest_sc_num], axis=1)


model_logit=sm.Logit(ytrain, xtrain).fit()
Optimization terminated successfully.
      Current function value: 0.234459
      Iterations 9

xtrain.isnull().sum()
Gender                                0
Education_Level                      0
Joining Designation                  0
Grade                                0
Quarterly Rating                     0
C10                                  0
C11                                  0
C12                                  0
C13                                  0
C14                                  0
C15                                  0
C16                                  0
C17                                  0
C18                                  0
C19                                  0
C2                                  0
C20                                  0
C21                                  0
C22                                  0

```

```
C23      0
C24      0
C25      0
C26      0
C27      0
C28      0
C29      0
C3       0
C4       0
C5       0
C6       0
C7       0
C8       0
C9       0
Age      0
Income   0
Total Business Value  0
dtype: int64
```

```
model_logit.summary()
```

```
<class 'statsmodels.iolib.summary.Summary'>
"""
```

Logit Regression Results

```
=====
Dep. Variable:          Attrition    No. Observations:
15283
Model:                  Logit        Df Residuals:
15247
Method:                 MLE          Df Model:
35
Date:                  Thu, 22 May 2025    Pseudo R-squ.:
0.1993
Time:                  16:46:13    Log-Likelihood:
-3583.2
converged:              True          LL-Null:
-4475.2
Covariance Type:        nonrobust      LLR p-value:
0.000
=====
```

```
=====
[0.025      0.975]
-----
Gender      0.0133      0.062      0.214      0.830
-0.109      0.135
Education_Level -0.0290      0.038     -0.762      0.446
```

-0.103	0.046				
Joining Designation		-0.0498	0.061	-0.818	0.413
-0.169	0.070				
Grade		-0.2705	0.064	-4.230	0.000
-0.396	-0.145				
Quarterly Rating		-1.4674	0.069	-21.367	0.000
-1.602	-1.333				
C10		-0.3670	0.208	-1.768	0.077
-0.774	0.040				
C11		-0.0880	0.226	-0.390	0.696
-0.530	0.354				
C12		-0.2025	0.203	-0.998	0.318
-0.600	0.195				
C13		-0.0937	0.209	-0.449	0.654
-0.503	0.316				
C14		-0.2617	0.206	-1.273	0.203
-0.665	0.141				
C15		-0.2082	0.192	-1.085	0.278
-0.584	0.168				
C16		-0.4263	0.218	-1.955	0.051
-0.854	0.001				
C17		0.0011	0.210	0.005	0.996
-0.410	0.412				
C18		-0.2960	0.226	-1.308	0.191
-0.739	0.148				
C19		-0.2285	0.225	-1.018	0.309
-0.669	0.211				
C2		-0.0691	0.207	-0.334	0.738
-0.475	0.336				
C20		-0.1539	0.173	-0.891	0.373
-0.493	0.185				
C21		-0.1959	0.215	-0.911	0.362
-0.618	0.226				
C22		-0.4444	0.211	-2.108	0.035
-0.858	-0.031				
C23		-0.1158	0.199	-0.583	0.560
-0.505	0.274				
C24		-0.2013	0.217	-0.928	0.354
-0.627	0.224				
C25		-0.2644	0.207	-1.279	0.201
-0.670	0.141				
C26		-0.3199	0.200	-1.599	0.110
-0.712	0.072				
C27		-0.2372	0.199	-1.189	0.234
-0.628	0.154				
C28		-0.1963	0.203	-0.965	0.334
-0.595	0.202				
C29		-0.3815	0.205	-1.860	0.063
-0.784	0.021				

C3		-0.3307	0.210	-1.575	0.115
-0.742	0.081				
C4		-0.1993	0.214	-0.931	0.352
-0.619	0.220				
C5		-0.3496	0.208	-1.678	0.093
-0.758	0.059				
C6		-0.0859	0.200	-0.429	0.668
-0.478	0.306				
C7		-0.3379	0.214	-1.576	0.115
-0.758	0.082				
C8		-0.1847	0.208	-0.890	0.374
-0.591	0.222				
C9		-0.1168	0.208	-0.562	0.574
-0.524	0.290				
Age		0.0970	0.034	2.869	0.004
0.031	0.163				
Income		0.0094	0.049	0.193	0.847
-0.086	0.105				
Total Business Value		-2.0511	0.174	-11.769	0.000
-2.393	-1.710				

```
=====
=====
"""
```

#So we see that features such as Gender, Income, Joining designation', and none of the cities are significant

#Only Age, Grade, Quarterly rating are significant and city C22

#Therefore we should ideally remove all the other features and put only these

xtest.dtypes

Gender	float64
Education_Level	int64
Joining Designation	int64
Grade	int64
Quarterly Rating	int64
C10	int32
C11	int32
C12	int32
C13	int32
C14	int32
C15	int32
C16	int32
C17	int32
C18	int32
C19	int32
C2	int32
C20	int32

```

C21          int32
C22          int32
C23          int32
C24          int32
C25          int32
C26          int32
C27          int32
C28          int32
C29          int32
C3           int32
C4           int32
C5           int32
C6           int32
C7           int32
C8           int32
C9           int32
Age          float64
Income       float64
Total Business Value float64
dtype: object

```

```
#xtest['tenure in days']=xtest['tenure in days'].astype('int')
```

```
ypred=model_logit.predict(xtest)
```

```
ypred_bin=[1 if i>=0.5 else 0 for i in ypred]
```

```
from sklearn.metrics import classification_report
```

```
print(classification_report(ytest, ypred_bin)) #Here we have a perfect recall for the class 0, however, we cannot predict the class 1 accurately at all
```

	precision	recall	f1-score	support
0	0.92	1.00	0.96	3517
1	0.00	0.00	0.00	304
accuracy			0.92	3821
macro avg	0.46	0.50	0.48	3821
weighted avg	0.85	0.92	0.88	3821

```
#trying with a different threshold
```

```
ypred_bin_2=[1 if i>=.2 else 0 for i in ypred] #So if we put a threshold of 0.1 we get a better result
```

```
print(classification_report(ytest, ypred_bin_2))
```

	precision	recall	f1-score	support
0	0.96	0.86	0.91	3517
1	0.27	0.58	0.37	304
accuracy			0.84	3821
macro avg	0.61	0.72	0.64	3821
weighted avg	0.90	0.84	0.87	3821

```
from sklearn.metrics import roc_curve, roc_auc_score
```


```
fpr, tpr, thresh=roc_curve(ytest, ypred)
```

```
plt.plot(fpr, tpr)
```

```
plt.ylabel('TPR')
```

```
plt.xlabel('FPR')
```

```
Text(0.5, 0, 'FPR')
```



```
print(roc_auc_score(ytest, ypred)) #So with logit itself the model  
gives an 82% score
```

0.8442382301004144

#Now I am trying with fewer features

```
x=x[['Age','Grade','Quarterly Rating']]
```

```
xtrain, xtest, ytrain, ytest=train_test_split(x, y, test_size=0.2,  
random_state=10)
```

```
model_logit=sm.Logit(ytrain, xtrain).fit()
```

```
Optimization terminated successfully.  
Current function value: 0.241661  
Iterations 8
```

```
model_logit.summary()
```

```
<class 'statsmodels.iolib.summary.Summary'>
```

```
"""
```

Logit Regression Results

```
=====
```

```
Dep. Variable:          Attrition    No. Observations:  
15283  
Model:                Logit        Df Residuals:  
15280  
Method:               MLE          Df Model:  
2  
Date:                 Thu, 22 May 2025    Pseudo R-squ.:  
0.1747  
Time:                 16:46:14    Log-Likelihood:  
-3693.3  
converged:            True        LL-Null:  
-4475.2  
Covariance Type:      nonrobust    LLR p-value:  
0.000
```

```
=====
```

```
=====
```

		coef	std err	z	P> z	
[0.025	0.975]					

Age		0.0268	0.003	8.659	0.000	
0.021	0.033					
Grade		-0.3080	0.033	-9.355	0.000	-
0.373	-0.243					
Quarterly Rating		-1.7667	0.065	-27.083	0.000	-
1.895	-1.639					

```
=====
=====
"""
```

```
from sklearn.metrics import recall_score, accuracy_score,
classification_report
```

```
ypred_logit=model_logit.predict(xtest)
```

```
ypred_logit_bin=[1 if i>=0.2 else 0 for i in ypred_logit]
```

```
print(classification_report(ytest, ypred_logit_bin))
```

	precision	recall	f1-score	support
0	0.94	0.91	0.92	3517
1	0.25	0.36	0.30	304
accuracy			0.86	3821
macro avg	0.60	0.63	0.61	3821
weighted avg	0.89	0.86	0.87	3821

```
from sklearn.tree import DecisionTreeClassifier
```

```
model_tree=DecisionTreeClassifier(random_state=10)
```

```
model_tree.fit(xtrain, ytrain)
```

```
DecisionTreeClassifier(random_state=10)
```

```
ypred_tree=model_tree.predict(xtest)
```

```
print(classification_report(ytest, ypred_tree))
```

	precision	recall	f1-score	support
0	0.92	1.00	0.96	3517
1	0.00	0.00	0.00	304
accuracy			0.92	3821
macro avg	0.46	0.50	0.48	3821
weighted avg	0.85	0.92	0.88	3821

```
ypred_train=model_tree.predict(xtrain)
```

```
print(classification_report(ytrain, ypred_train))
```

#There is no difference between training and test results, so there is no need for regularization or bagging

	precision	recall	f1-score	support
--	-----------	--------	----------	---------

0	0.91	1.00	0.96	13971
1	0.70	0.01	0.02	1312
accuracy			0.91	15283
macro avg	0.81	0.51	0.49	15283
weighted avg	0.90	0.91	0.88	15283

#Hence trying to adjust max depth

`model_tree_1=DecisionTreeClassifier(max_depth=5)` *#Trying with a small depth, and that appears to make the class 1 recall completely bad*

`model_tree_1.fit(xtrain, ytrain)`

`DecisionTreeClassifier(max_depth=5)`

`ypred_tree_1=model_tree_1.predict(xtest)`

`print(classification_report(ytest, ypred_tree_1))`

	precision	recall	f1-score	support
0	0.92	1.00	0.96	3517
1	0.00	0.00	0.00	304
accuracy			0.92	3821
macro avg	0.46	0.50	0.48	3821
weighted avg	0.85	0.92	0.88	3821

`model_tree_2=DecisionTreeClassifier(max_depth=30)` *#increasing the depth to see if this gives a better recall for class 1*

`model_tree_2.fit(xtrain, ytrain)`

`DecisionTreeClassifier(max_depth=30)`

`ypred_tree_2=model_tree_2.predict(xtest)`

`print(classification_report(ytest, ypred_tree_2))` *#it still shows very poor results for class 1*

	precision	recall	f1-score	support
0	0.92	1.00	0.96	3517
1	0.00	0.00	0.00	304
accuracy			0.92	3821
macro avg	0.46	0.50	0.48	3821
weighted avg	0.85	0.92	0.88	3821

#This shows the model needs a very high depth to start classifying the results of class 1 correctly

#Hence trying GradientBoosting

```
from sklearn.ensemble import GradientBoostingClassifier
```

```
model_gb=GradientBoostingClassifier(random_state=10)
```

```
model_gb.fit(xtrain, ytrain)
```

```
GradientBoostingClassifier(random_state=10)
```

```
ypred_gb_0=model_gb.predict(xtest)
```

```
print(classification_report(ytest, ypred_gb_0)) #Gradient Boosting does not improve it directly, we see we are still getting only 1% recall of class 1
```

	precision	recall	f1-score	support
0	0.92	1.00	0.96	3517
1	0.00	0.00	0.00	304
accuracy			0.92	3821
macro avg	0.46	0.50	0.48	3821
weighted avg	0.85	0.92	0.88	3821

```
ypred_gb_train=model_gb.predict(xtrain)
```

```
print(classification_report(ytrain, ypred_gb_train))
```

#So doing gradient boosting is not helping here, it is underfitting because the values of the 1 class are not being predicted correctly

	precision	recall	f1-score	support
0	0.91	1.00	0.96	13971
1	0.67	0.01	0.02	1312
accuracy			0.91	15283
macro avg	0.79	0.50	0.49	15283
weighted avg	0.89	0.91	0.87	15283

#Hence trying to increase the samples using SMOTE

```
import imblearn
```

```
from imblearn.over_sampling import SMOTE
```

```
smote=SMOTE()
```

```
xtrain_sm, ytrain_sm=smote.fit_resample(xtrain, ytrain)

model_logit_2=sm.Logit(ytrain_sm, xtrain_sm).fit() #trying logit with SMOTE

Optimization terminated successfully.
      Current function value: 0.512970
      Iterations 7

model_logit_2.summary()

<class 'statsmodels.iolib.summary.Summary'>
"""
                                Logit Regression Results

=====
=====
Dep. Variable:                  Attrition   No. Observations:
27942
Model:                          Logit      Df Residuals:
27939
Method:                         MLE        Df Model:
2
Date:                           Thu, 22 May 2025   Pseudo R-squ.:
0.2599
Time:                           16:46:16   Log-Likelihood:
-14333.
converged:                       True      LL-Null:
-19368.
Covariance Type:                nonrobust   LLR p-value:
0.000
=====
=====
                                coef      std err          z      P>|z|
-----
[0.025      0.975]
-----
Age                                0.1002      0.002      64.003      0.000
0.097      0.103
Grade                             -0.2930      0.016     -18.716      0.000      -
0.324     -0.262
Quarterly Rating                  -2.0062      0.029     -68.346      0.000      -
2.064     -1.949
=====
=====
"""

ypred_logit_2=model_logit_2.predict(xtest)

ypred_bin=[1 if i>=0.5 else 0 for i in ypred_logit_2]
```

```
print(classification_report(ytest, ypred_bin)) #so using smote with
logit clearly helps
```

	precision	recall	f1-score	support
0	0.98	0.66	0.79	3517
1	0.18	0.86	0.29	304
accuracy			0.67	3821
macro avg	0.58	0.76	0.54	3821
weighted avg	0.92	0.67	0.75	3821

```
ypred_bin2=[1 if i>=0.1 else 0 for i in ypred_logit_2] #by reducing
the threshold, the overall accuracy has dipped
```

*#This gives me a clue that I should try stacking models and put a
logistic regression at the end and lower the threshold*

```
print(classification_report(ytest, ypred_bin2)) #This now gives the
highest recall for class 1, this is the overall best result
```

	precision	recall	f1-score	support
0	1.00	0.32	0.48	3517
1	0.11	0.99	0.20	304
accuracy			0.37	3821
macro avg	0.55	0.65	0.34	3821
weighted avg	0.93	0.37	0.46	3821

```
model_tree=DecisionTreeClassifier(random_state=10)
```

```
model_tree.fit(xtrain_sm, ytrain_sm)
```

```
DecisionTreeClassifier(random_state=10)
```

```
ypred_tree_sm=model_tree.predict(xtest)
```

```
print(classification_report(ytest, ypred_tree_sm))
```

	precision	recall	f1-score	support
0	0.98	0.68	0.80	3517
1	0.18	0.84	0.30	304
accuracy			0.69	3821
macro avg	0.58	0.76	0.55	3821
weighted avg	0.92	0.69	0.76	3821

```

model_gb=GradientBoostingClassifier(random_state=10)
model_gb.fit(xtrain_sm,ytrain_sm)
GradientBoostingClassifier(random_state=10)
ypred_gb2=model_gb.predict(xtest)
print(classification_report(ytest, ypred_gb2))
#using a combination of smote with gradient boosting classifier has yielded the best results so far

```

	precision	recall	f1-score	support
0	0.98	0.68	0.80	3517
1	0.18	0.85	0.30	304
accuracy			0.69	3821
macro avg	0.58	0.76	0.55	3821
weighted avg	0.92	0.69	0.76	3821

#Can I increase it further and ensure that the recall of both classes is coming out well

#So I am trying to decrease the number of estimators maybe because with more estimators in parallel, it is not giving a great accuracy

```

model_gb3=GradientBoostingClassifier(n_estimators=20, random_state=20)
model_gb3.fit(xtrain_sm, ytrain_sm)
GradientBoostingClassifier(n_estimators=20, random_state=20)
ypred_gb3=model_gb3.predict(xtest)
print(classification_report(ytest, ypred_gb3))

```

	precision	recall	f1-score	support
0	0.98	0.65	0.78	3517
1	0.18	0.88	0.30	304
accuracy			0.67	3821
macro avg	0.58	0.77	0.54	3821
weighted avg	0.92	0.67	0.75	3821

#I then tried reducing learning rate to see if I can increase the recall of the class 1 further without sacrificing overall accuracy too much

```
model_gb4=GradientBoostingClassifier(n_estimators=10,
learning_rate=.05, random_state=10)
```

```
model_gb4.fit(xtrain_sm, ytrain_sm)
```

```
GradientBoostingClassifier(learning_rate=0.05, n_estimators=10,
random_state=10)
```

```
ypred_gb4=model_gb4.predict(xtest)
```

```
print(classification_report(ytest, ypred_gb4))
```

	precision	recall	f1-score	support
0	0.98	0.69	0.81	3517
1	0.19	0.84	0.31	304
accuracy			0.70	3821
macro avg	0.58	0.76	0.56	3821
weighted avg	0.92	0.70	0.77	3821

#then i tried a max depth of 3 and it worked better, reducing depth below that is not improving anything

```
model_gb5=GradientBoostingClassifier(n_estimators=10,
learning_rate=0.05, max_depth=3, random_state=10)
```

```
model_gb5.fit(xtrain_sm, ytrain_sm)
```

```
GradientBoostingClassifier(learning_rate=0.05, n_estimators=10,
random_state=10)
```

```
ypred_gb5=model_gb5.predict(xtest)
```

```
print(classification_report(ytest, ypred_gb5))
```


	precision	recall	f1-score	support
0	0.98	0.69	0.81	3517
1	0.19	0.84	0.31	304
accuracy			0.70	3821
macro avg	0.58	0.76	0.56	3821
weighted avg	0.92	0.70	0.77	3821

```
fpr, tpr, thresh=roc_curve(ytest, ypred_gb5)
```

```
plt.plot(fpr, tpr)
```

```
plt.title('ROC AUC curve for the gradient boosting classifier')
```

```
Text(0.5, 1.0, 'ROC AUC curve for the gradient boosting classifier')
```



```
print(roc_auc_score(ytest, ypred_gb5))
```

```
0.7639585172769854
```

*#i understand that there is class imbalance, and to address that, one of the ways is to use stratified kfold sampling
#So I will use stratified kfold sampling on top of the smote data to see if it improves results*

```
from sklearn.model_selection import StratifiedKFold
```

```
skf=StratifiedKFold(n_splits=10,shuffle=True, random_state=10)
```

```
from sklearn.metrics import recall_score, accuracy_score
```

```
model_gb5=GradientBoostingClassifier(n_estimators=20,  
learning_rate=0.05, max_depth=3,random_state=10)
```

```
model_recall_score=[]
```

```
model_acc_score=[]
```

```
for train_index, test_index in skf.split(x,y):  
    xtrain_fold, xtest_fold=x.iloc[train_index], x.iloc[test_index]
```

```

ytrain_fold, ytest_fold=y.iloc[train_index], y.iloc[test_index]
xtrain_sm_fold, ytrain_sm_fold=smote.fit_resample(xtrain_fold,
ytrain_fold)
model_gb5.fit(xtrain_sm_fold, ytrain_sm_fold)
ypred_gb5=model_gb5.predict(xtest_fold)
model_recall_score.append(recall_score(ytest_fold, ypred_gb5))
model_acc_score.append(accuracy_score(ytest_fold, ypred_gb5))

```

model_recall_score

```

[0.8580246913580247,
 0.8333333333333334,
 0.8518518518518519,
 0.8333333333333334,
 0.8260869565217391,
 0.84472049689441,
 0.8881987577639752,
 0.8571428571428571,
 0.8209876543209876,
 0.845679012345679]

```

np.mean(model_recall_score) #this is still not giving great recall even though for some row sample sets it is crossing 85%

0.8459358944866191

```
import xgboost
```

```
from xgboost import XGBClassifier
```

model_xgb=XGBClassifier(random_state=10) #The xgboost classifier does not provide a good recall here

```

model_recall_score_xgb=[]
for train_index, test_index in skf.split(x,y):
    xtrain_fold, xtest_fold=x.iloc[train_index], x.iloc[test_index]
    ytrain_fold, ytest_fold=y.iloc[train_index], y.iloc[test_index]
    xtrain_sm_fold, ytrain_sm_fold=smote.fit_resample(xtrain_fold,
ytrain_fold)
    model_xgb.fit(xtrain_sm_fold, ytrain_sm_fold)
    ypred_xgb_kf=model_xgb.predict(xtest_fold)
    model_recall_score_xgb.append(recall_score(ytest_fold,
ypred_xgb_kf))

```

model_recall_score_xgb

```

[0.8580246913580247,
 0.8641975308641975,
 0.8395061728395061,
 0.8148148148148148,
 0.8385093167701864,

```



```
0.8509316770186336,  
0.8757763975155279,  
0.8385093167701864,  
0.8148148148148148,  
0.845679012345679]
```

```
from sklearn.ensemble import RandomForestClassifier  
model_RF=RandomForestClassifier(max_depth=3, n_estimators=20)  
  
model_recall_score_rf=[]  
for train_index, test_index in skf.split(x,y):  
    xtrain_fold, xtest_fold=x.iloc[train_index], x.iloc[test_index]  
    ytrain_fold, ytest_fold=y.iloc[train_index], y.iloc[test_index]  
    xtrain_sm_fold, ytrain_sm_fold=smote.fit_resample(xtrain_fold,  
ytrain_fold)  
    model_RF.fit(xtrain_sm_fold, ytrain_sm_fold)  
    ypred_rf_kf=model_RF.predict(xtest_fold)  
    model_recall_score_rf.append(recall_score(ytest_fold,  
ypred_rf_kf))
```

```
model_recall_score_rf
```

```
[0.8580246913580247,  
0.8333333333333334,  
0.8518518518518519,  
0.8333333333333334,  
0.8260869565217391,  
0.84472049689441,  
0.8819875776397516,  
0.84472049689441,  
0.8024691358024691,  
0.845679012345679]
```

```
model_RF.fit(xtrain_sm, ytrain_sm)
```

```
RandomForestClassifier(max_depth=3, n_estimators=20)
```

```
ypred_rf=model_RF.predict(xtest)
```

```
print(classification_report(ytest, ypred_rf)) #this works because we  
have only 3 features
```

	precision	recall	f1-score	support
0	0.98	0.69	0.81	3517
1	0.19	0.84	0.31	304
accuracy			0.70	3821

macro avg	0.58	0.76	0.56	3821
weighted avg	0.92	0.70	0.77	3821

```
#pip install xgboost
```

```
from xgboost import XGBClassifier
```

```
model_xgb=XGBClassifier()
```

```
model_xgb.fit(xtrain, ytrain)
```

```
XGBClassifier(base_score=None, booster=None, callbacks=None,
               colsample_bylevel=None, colsample_bynode=None,
               colsample_bytree=None, device=None,
               early_stopping_rounds=None,
               enable_categorical=False, eval_metric=None,
               feature_types=None,
               feature_weights=None, gamma=None, grow_policy=None,
               importance_type=None, interaction_constraints=None,
               learning_rate=None, max_bin=None,
               max_cat_threshold=None,
               max_cat_to_onehot=None, max_delta_step=None,
               max_depth=None,
               max_leaves=None, min_child_weight=None, missing=nan,
               monotone_constraints=None, multi_strategy=None,
               n_estimators=None,
               n_jobs=None, num_parallel_tree=None, ...)
```

```
ypred_xgb_1=model_xgb.predict(xtest)
```

```
print(classification_report(ytest, ypred_xgb_1))
```

	precision	recall	f1-score	support
0	0.92	1.00	0.96	3517
1	0.00	0.00	0.00	304
accuracy			0.92	3821
macro avg	0.46	0.50	0.48	3821
weighted avg	0.85	0.92	0.88	3821

```
model_xgb2=XGBClassifier()
```

```
model_xgb2.fit(xtrain_sm, ytrain_sm)
```

```
XGBClassifier(base_score=None, booster=None, callbacks=None,
               colsample_bylevel=None, colsample_bynode=None,
               colsample_bytree=None, device=None,
```

```

early_stopping_rounds=None,
        enable_categorical=False, eval_metric=None,
feature_types=None,
        feature_weights=None, gamma=None, grow_policy=None,
        importance_type=None, interaction_constraints=None,
        learning_rate=None, max_bin=None,
max_cat_threshold=None,
        max_cat_to_onehot=None, max_delta_step=None,
max_depth=None,
        max_leaves=None, min_child_weight=None, missing=nan,
        monotone_constraints=None, multi_strategy=None,
n_estimators=None,
        n_jobs=None, num_parallel_tree=None, ...)

ypred_xgb2=model_xgb2.predict(xtest)

print(classification_report(ytest, ypred_xgb2)) #We see that xgboost
does not address my recall problem with class 1, it only increases
overall accuracy

```

	precision	recall	f1-score	support
0	0.98	0.67	0.80	3517
1	0.18	0.85	0.30	304
accuracy			0.69	3821
macro avg	0.58	0.76	0.55	3821
weighted avg	0.92	0.69	0.76	3821

```

from sklearn.ensemble import AdaBoostClassifier

model_adb=AdaBoostClassifier()

model_adb.fit(xtrain_sm, ytrain_sm)

```

E:\python\Lib\site-packages\sklearn\ensemble_weight_boosting.py:527: FutureWarning: The SAMME.R algorithm (the default) is deprecated and will be removed in 1.6. Use the SAMME algorithm to circumvent this warning.

```

    warnings.warn(

AdaBoostClassifier()

ypred_adb=model_adb.predict(xtest)

print(classification_report(ytest, ypred_adb))

```

	precision	recall	f1-score	support
--	-----------	--------	----------	---------

	0	0.98	0.69	0.81	3517
	1	0.19	0.84	0.31	304
accuracy				0.70	3821
macro avg		0.58	0.76	0.56	3821
weighted avg		0.92	0.70	0.77	3821

```
model_adb=AdaBoostClassifier(n_estimators=10, learning_rate=.05,
random_state=10)
```

```
model_adb.fit(xtrain_sm, ytrain_sm)
```

E:\python\Lib\site-packages\sklearn\ensemble_weight_boosting.py:527: FutureWarning: The SAMME.R algorithm (the default) is deprecated and will be removed in 1.6. Use the SAMME algorithm to circumvent this warning.

```
warnings.warn(
```

```
AdaBoostClassifier(learning_rate=0.05, n_estimators=10,
random_state=10)
```

```
ypred_adb_2=model_adb.predict(xtest)
```

```
print(classification_report(ytest, ypred_adb_2))
```

		precision	recall	f1-score	support
	0	0.99	0.64	0.78	3517
	1	0.18	0.89	0.30	304
accuracy				0.66	3821
macro avg		0.58	0.77	0.54	3821
weighted avg		0.92	0.66	0.74	3821

```
model_recall_score_adb=[]
```

```
model_acc_score=[]
```

```
for train_index, test_index in skf.split(x,y):
    xtrain_fold, xtest_fold=x.iloc[train_index], x.iloc[test_index]
    ytrain_fold, ytest_fold=y.iloc[train_index], y.iloc[test_index]
    xtrain_sm_fold, ytrain_sm_fold=smote.fit_resample(xtrain_fold,
ytrain_fold)
    model_adb.fit(xtrain_sm_fold, ytrain_sm_fold)
    ypred_adb=model_adb.predict(xtest_fold)
    model_recall_score.append(recall_score(ytest_fold, ypred_adb))
    model_acc_score.append(accuracy_score(ytest_fold, ypred_adb))
```

E:\python\Lib\site-packages\sklearn\ensemble_weight_boosting.py:527: FutureWarning: The SAMME.R algorithm (the default) is deprecated and will be removed in 1.6. Use the SAMME algorithm to circumvent this warning.

[illegible]

```
random_state=10))],  
        final_estimator=LogisticRegression())
```

```
stk_clf.fit(xtrain_sm, ytrain_sm)
```

```
E:\python\Lib\site-packages\sklearn\ensemble\_weight_boosting.py:527:
FutureWarning: The SAMME.R algorithm (the default) is deprecated and
will be removed in 1.6. Use the SAMME algorithm to circumvent this
warning.
```

```
    warnings.warn(
```

```
E:\python\Lib\site-packages\sklearn\ensemble\_weight_boosting.py:527:
FutureWarning: The SAMME.R algorithm (the default) is deprecated and
will be removed in 1.6. Use the SAMME algorithm to circumvent this
warning.
```

```
    warnings.warn(
```

```
E:\python\Lib\site-packages\sklearn\ensemble\_weight_boosting.py:527:
FutureWarning: The SAMME.R algorithm (the default) is deprecated and
will be removed in 1.6. Use the SAMME algorithm to circumvent this
warning.
```

```
    warnings.warn(
```

```
E:\python\Lib\site-packages\sklearn\ensemble\_weight_boosting.py:527:
FutureWarning: The SAMME.R algorithm (the default) is deprecated and
will be removed in 1.6. Use the SAMME algorithm to circumvent this
warning.
```

```
    warnings.warn(
```

```
E:\python\Lib\site-packages\sklearn\ensemble\_weight_boosting.py:527:
FutureWarning: The SAMME.R algorithm (the default) is deprecated and
will be removed in 1.6. Use the SAMME algorithm to circumvent this
warning.
```

```
    warnings.warn(
```

```
E:\python\Lib\site-packages\sklearn\ensemble\_weight_boosting.py:527:
FutureWarning: The SAMME.R algorithm (the default) is deprecated and
will be removed in 1.6. Use the SAMME algorithm to circumvent this
warning.
```

```
    warnings.warn(
```

```
StackingClassifier(estimators=[('ADB',
                                AdaBoostClassifier(learning_rate=0.05,
                                                       n_estimators=10,
                                                       random_state=10)),
                                ('GBM',
```

```
GradientBoostingClassifier(learning_rate=0.05,
```

```
n_estimators=20,
```

```
random_state=10))],
```

```
        final_estimator=LogisticRegression())
```

```
ypred_stk=stk_clf.predict(xtest)
```

```
print(classification_report(ytest, ypred_stk))
```

	precision	recall	f1-score	support
0	0.98	0.69	0.81	3517
1	0.19	0.84	0.31	304
accuracy			0.70	3821
macro avg	0.58	0.76	0.56	3821
weighted avg	0.92	0.70	0.77	3821

```
ypred_stk_proba=stk_clf.predict_proba(xtest)
ypred_stk_proba=ypred_stk_proba[:,1] #take only probability of 1 class
ypred_stk_bin=[1 if i>=0.1 else 0 for i in ypred_stk_proba]
print(classification_report(ytest, ypred_stk_bin)) #this is the second best result
```

	precision	recall	f1-score	support
0	0.99	0.46	0.63	3517
1	0.13	0.95	0.23	304
accuracy			0.50	3821
macro avg	0.56	0.71	0.43	3821
weighted avg	0.92	0.50	0.60	3821

```
pip install LightGBM
```

Requirement already satisfied: LightGBM in e:\python\lib\site-packages (4.6.0)

Requirement already satisfied: numpy>=1.17.0 in e:\python\lib\site-packages (from LightGBM) (1.26.4)

Requirement already satisfied: scipy in e:\python\lib\site-packages (from LightGBM) (1.13.1)

Note: you may need to restart the kernel to use updated packages.

```
import lightgbm as lgb
```

```
model_lgb=lgb.LGBMClassifier(n_estimators=10, learning_rate=.05)
```

```
model_lgb.fit(xtrain_sm, ytrain_sm)
```

[LightGBM] [Warning] Found whitespace in feature_names, replace with underlines

[LightGBM] [Info] Number of positive: 13971, number of negative: 13971

[LightGBM] [Info] Auto-choosing col-wise multi-threading, the overhead of testing was 0.000904 seconds.

You can set `force_col_wise=true` to remove the overhead.


```
[LightGBM] [Info] Total Bins 208
[LightGBM] [Info] Number of data points in the train set: 27942,
number of used features: 3
[LightGBM] [Info] [binary:BoostFromScore]: pavg=0.500000 ->
initscore=0.000000
```

```
LGBMClassifier(learning_rate=0.05, n_estimators=10)
```

```
ypred_lgb=model_lgb.predict(xtest)
```

```
print(classification_report(ytest, ypred_lgb))
```

	precision	recall	f1-score	support
0	0.98	0.68	0.80	3517
1	0.18	0.85	0.30	304
accuracy			0.69	3821
macro avg	0.58	0.76	0.55	3821
weighted avg	0.92	0.69	0.76	3821

```
model_stk_2=StackingClassifier(estimators=[('lgb', model_lgb),
('ADB',model_adb), ('GBM', model_gb5)],
final_estimator=LogisticRegression())
```

```
model_stk_2.fit(xtrain_sm, ytrain_sm)
```

```
[LightGBM] [Warning] Found whitespace in feature_names, replace with
underscores
[LightGBM] [Info] Number of positive: 13971, number of negative: 13971
[LightGBM] [Info] Auto-choosing row-wise multi-threading, the overhead
of testing was 0.000163 seconds.
You can set `force_row_wise=true` to remove the overhead.
And if memory is not enough, you can set `force_col_wise=true`.
[LightGBM] [Info] Total Bins 208
[LightGBM] [Info] Number of data points in the train set: 27942,
number of used features: 3
[LightGBM] [Info] [binary:BoostFromScore]: pavg=0.500000 ->
initscore=0.000000
```

```
E:\python\Lib\site-packages\sklearn\ensemble\_weight_boosting.py:527:
FutureWarning: The SAMME.R algorithm (the default) is deprecated and
will be removed in 1.6. Use the SAMME algorithm to circumvent this
warning.
```

```
warnings.warn(
```

```
[LightGBM] [Warning] Found whitespace in feature_names, replace with
underscores
[LightGBM] [Info] Number of positive: 11177, number of negative: 11176
[LightGBM] [Info] Auto-choosing col-wise multi-threading, the overhead
of testing was 0.011635 seconds.
```

You can set `force_col_wise=true` to remove the overhead.
[LightGBM] [Info] Total Bins 218
[LightGBM] [Info] Number of data points in the train set: 22353,
number of used features: 3
[LightGBM] [Info] [binary:BoostFromScore]: pavg=0.500022 ->
initscore=0.000089
[LightGBM] [Info] Start training from score 0.000089
[LightGBM] [Warning] Found whitespace in feature_names, replace with
underscores
[LightGBM] [Info] Number of positive: 11176, number of negative: 11177
[LightGBM] [Info] Auto-choosing row-wise multi-threading, the overhead
of testing was 0.000248 seconds.
You can set `force_row_wise=true` to remove the overhead.
And if memory is not enough, you can set `force_col_wise=true`.
[LightGBM] [Info] Total Bins 212
[LightGBM] [Info] Number of data points in the train set: 22353,
number of used features: 3
[LightGBM] [Info] [binary:BoostFromScore]: pavg=0.499978 ->
initscore=-0.000089
[LightGBM] [Info] Start training from score -0.000089
[LightGBM] [Warning] Found whitespace in feature_names, replace with
underscores
[LightGBM] [Info] Number of positive: 11177, number of negative: 11177
[LightGBM] [Info] Auto-choosing row-wise multi-threading, the overhead
of testing was 0.000242 seconds.
You can set `force_row_wise=true` to remove the overhead.
And if memory is not enough, you can set `force_col_wise=true`.
[LightGBM] [Info] Total Bins 212
[LightGBM] [Info] Number of data points in the train set: 22354,
number of used features: 3
[LightGBM] [Info] [binary:BoostFromScore]: pavg=0.500000 ->
initscore=0.000000
[LightGBM] [Warning] Found whitespace in feature_names, replace with
underscores
[LightGBM] [Info] Number of positive: 11177, number of negative: 11177
[LightGBM] [Info] Auto-choosing row-wise multi-threading, the overhead
of testing was 0.000245 seconds.
You can set `force_row_wise=true` to remove the overhead.
And if memory is not enough, you can set `force_col_wise=true`.
[LightGBM] [Info] Total Bins 205
[LightGBM] [Info] Number of data points in the train set: 22354,
number of used features: 3
[LightGBM] [Info] [binary:BoostFromScore]: pavg=0.500000 ->
initscore=0.000000
[LightGBM] [Warning] Found whitespace in feature_names, replace with
underscores
[LightGBM] [Info] Number of positive: 11177, number of negative: 11177
[LightGBM] [Info] Auto-choosing row-wise multi-threading, the overhead
of testing was 0.000205 seconds.

You can set `force_row_wise=True` to remove the overhead.
And if memory is not enough, you can set `force_col_wise=True`.
[LightGBM] [Info] Total Bins 207
[LightGBM] [Info] Number of data points in the train set: 22354,
number of used features: 3
[LightGBM] [Info] [binary:BoostFromScore]: pavg=0.500000 ->
initscore=0.000000

E:\python\Lib\site-packages\sklearn\ensemble_weight_boosting.py:527:
FutureWarning: The SAMME.R algorithm (the default) is deprecated and
will be removed in 1.6. Use the SAMME algorithm to circumvent this
warning.

```
warnings.warn(
```

E:\python\Lib\site-packages\sklearn\ensemble_weight_boosting.py:527:
FutureWarning: The SAMME.R algorithm (the default) is deprecated and
will be removed in 1.6. Use the SAMME algorithm to circumvent this
warning.

```
warnings.warn(
```

E:\python\Lib\site-packages\sklearn\ensemble_weight_boosting.py:527:
FutureWarning: The SAMME.R algorithm (the default) is deprecated and
will be removed in 1.6. Use the SAMME algorithm to circumvent this
warning.

```
warnings.warn(
```

E:\python\Lib\site-packages\sklearn\ensemble_weight_boosting.py:527:
FutureWarning: The SAMME.R algorithm (the default) is deprecated and
will be removed in 1.6. Use the SAMME algorithm to circumvent this
warning.

```
warnings.warn(
```

E:\python\Lib\site-packages\sklearn\ensemble_weight_boosting.py:527:
FutureWarning: The SAMME.R algorithm (the default) is deprecated and
will be removed in 1.6. Use the SAMME algorithm to circumvent this
warning.

```
warnings.warn(
```

```
StackingClassifier(estimators=[('lgb',  
                                LGBMClassifier(learning_rate=0.05,  
                                                n_estimators=10)),  
                                ('ADB',  
                                 AdaBoostClassifier(learning_rate=0.05,  
                                                    n_estimators=10,  
                                                    random_state=10)),  
                                ('GBM',
```

```
GradientBoostingClassifier(learning_rate=0.05,
```

```
n_estimators=20,
```

```
random_state=10))],
```

```
    final_estimator=LogisticRegression())
```

```
ypred_stk_2=model_stk_2.predict(xtest)
```

```
print(classification_report(ytest, ypred_stk_2))
```

	precision	recall	f1-score	support
0	0.98	0.69	0.81	3517
1	0.19	0.84	0.31	304
accuracy			0.70	3821
macro avg	0.58	0.76	0.56	3821
weighted avg	0.92	0.70	0.77	3821

```
ypred_proba_stk_2=model_stk_2.predict_proba(xtest)[: ,1]
```

```
ypred_proba_stk_2
```

```
array([0.06812235, 0.72639533, 0.09733343, ..., 0.67249155,  
0.06733472,  
0.05501743])
```

```
ypred_bin_stk_2=[1 if i>=.1 else 0 for i in ypred_proba_stk_2]
```

```
print(classification_report(ytest, ypred_bin_stk_2)) #so this gives me  
a slightly better overall accuracy
```

	precision	recall	f1-score	support
0	0.99	0.53	0.69	3517
1	0.14	0.92	0.25	304
accuracy			0.56	3821
macro avg	0.57	0.73	0.47	3821
weighted avg	0.92	0.56	0.65	3821

```
pip install catboost
```

```
Requirement already satisfied: catboost in e:\python\lib\site-packages  
(1.2.8)
```

```
Requirement already satisfied: graphviz in e:\python\lib\site-packages  
(from catboost) (0.20.3)
```

```
Requirement already satisfied: matplotlib in e:\python\lib\site-  
packages (from catboost) (3.9.2)
```

```
Requirement already satisfied: numpy<3.0,>=1.16.0 in e:\python\lib\  
site-packages (from catboost) (1.26.4)
```

```
Requirement already satisfied: pandas>=0.24 in e:\python\lib\site-  
packages (from catboost) (2.2.2)
```

```
Requirement already satisfied: scipy in e:\python\lib\site-packages
```

```

(from catboost) (1.13.1)
Requirement already satisfied: plotly in e:\python\lib\site-packages
(from catboost) (5.24.1)
Requirement already satisfied: six in e:\python\lib\site-packages
(from catboost) (1.16.0)
Requirement already satisfied: python-dateutil>=2.8.2 in e:\python\
lib\site-packages (from pandas>=0.24->catboost) (2.9.0.post0)
Requirement already satisfied: pytz>=2020.1 in e:\python\lib\site-
packages (from pandas>=0.24->catboost) (2024.1)
Requirement already satisfied: tzdata>=2022.7 in e:\python\lib\site-
packages (from pandas>=0.24->catboost) (2023.3)
Requirement already satisfied: contourpy>=1.0.1 in e:\python\lib\site-
packages (from matplotlib->catboost) (1.2.0)
Requirement already satisfied: cycler>=0.10 in e:\python\lib\site-
packages (from matplotlib->catboost) (0.11.0)
Requirement already satisfied: fonttools>=4.22.0 in e:\python\lib\
site-packages (from matplotlib->catboost) (4.51.0)
Requirement already satisfied: kiwisolver>=1.3.1 in e:\python\lib\
site-packages (from matplotlib->catboost) (1.4.4)
Requirement already satisfied: packaging>=20.0 in e:\python\lib\site-
packages (from matplotlib->catboost) (24.1)
Requirement already satisfied: pillow>=8 in e:\python\lib\site-
packages (from matplotlib->catboost) (10.4.0)
Requirement already satisfied: pyparsing>=2.3.1 in e:\python\lib\site-
packages (from matplotlib->catboost) (3.1.2)
Requirement already satisfied: tenacity>=6.2.0 in e:\python\lib\site-
packages (from plotly->catboost) (8.2.3)
Note: you may need to restart the kernel to use updated packages.

```

```
import catboost
```

```
from catboost import CatBoostClassifier, Pool
```

```
model_catboost=CatBoostClassifier(learning_rate=0.01, depth=3,
iterations=2) #this works well when I simplify it by reducing the
depth and complexity
```

```
model_catboost.fit(xtrain_sm, ytrain_sm)
```

```
0:   learn: 0.6895776 total: 4.05ms   remaining: 4.05ms
1:   learn: 0.6859223 total: 8.77ms   remaining: 0us
```

```
<catboost.core.CatBoostClassifier at 0x1659b738ad0>
```

```
ypred_cat=model_catboost.predict(xtest)
```

```
print(classification_report(ytest, ypred_cat))
```

	precision	recall	f1-score	support
0	0.99	0.64	0.78	3517

1	0.18	0.89	0.30	304
accuracy			0.66	3821
macro avg	0.58	0.77	0.54	3821
weighted avg	0.92	0.66	0.74	3821

```
model_stk_3=StackingClassifier(estimators=[('catboost',model_catboost),
,('ADB',model_adb),('GBM',model_gb5)],
final_estimator=LogisticRegression())
```

```
model_stk_3.fit(xtrain_sm, ytrain_sm)
```

```
0:   learn: 0.6895776 total: 4.76ms   remaining: 4.76ms
1:   learn: 0.6859223 total: 9.39ms   remaining: 0us
```

```
E:\python\Lib\site-packages\sklearn\ensemble\_weight_boosting.py:527:
FutureWarning: The SAMME.R algorithm (the default) is deprecated and
will be removed in 1.6. Use the SAMME algorithm to circumvent this
warning.
```

```
warnings.warn(
```

```
0:   learn: 0.6894442 total: 4.89ms   remaining: 4.89ms
1:   learn: 0.6856558 total: 9.95ms   remaining: 0us
0:   learn: 0.6896024 total: 5.25ms   remaining: 5.25ms
1:   learn: 0.6859776 total: 10.1ms   remaining: 0us
0:   learn: 0.6896491 total: 4.66ms   remaining: 4.66ms
1:   learn: 0.6860893 total: 10.3ms   remaining: 0us
0:   learn: 0.6895973 total: 7.5ms    remaining: 7.5ms
1:   learn: 0.6859621 total: 16.3ms   remaining: 0us
0:   learn: 0.6895842 total: 4.9ms    remaining: 4.9ms
1:   learn: 0.6859574 total: 9.64ms   remaining: 0us
```

```
E:\python\Lib\site-packages\sklearn\ensemble\_weight_boosting.py:527:
FutureWarning: The SAMME.R algorithm (the default) is deprecated and
will be removed in 1.6. Use the SAMME algorithm to circumvent this
warning.
```

```
warnings.warn(
```

```
E:\python\Lib\site-packages\sklearn\ensemble\_weight_boosting.py:527:
FutureWarning: The SAMME.R algorithm (the default) is deprecated and
will be removed in 1.6. Use the SAMME algorithm to circumvent this
warning.
```

```
warnings.warn(
```

```
E:\python\Lib\site-packages\sklearn\ensemble\_weight_boosting.py:527:
FutureWarning: The SAMME.R algorithm (the default) is deprecated and
will be removed in 1.6. Use the SAMME algorithm to circumvent this
warning.
```

```
warnings.warn(
```

```
E:\python\Lib\site-packages\sklearn\ensemble\_weight_boosting.py:527:
```

```
FutureWarning: The SAMME.R algorithm (the default) is deprecated and
will be removed in 1.6. Use the SAMME algorithm to circumvent this
warning.
```

```
warnings.warn(
E:\python\Lib\site-packages\sklearn\ensemble\_weight_boosting.py:527:
FutureWarning: The SAMME.R algorithm (the default) is deprecated and
will be removed in 1.6. Use the SAMME algorithm to circumvent this
warning.
```

```
warnings.warn(
StackingClassifier(estimators=[('catboost',
                                <catboost.core.CatBoostClassifier
object at 0x000001659B738AD0>),
                                ('ADB',
                                 AdaBoostClassifier(learning_rate=0.05,
                                                       n_estimators=10,
                                                       random_state=10))),
                                ('GBM',
```

```
GradientBoostingClassifier(learning_rate=0.05,
n_estimators=20,
random_state=10))],
                        final_estimator=LogisticRegression())
```

```
ypred_stk_3=model_stk_3.predict(xtest)
```

```
print(classification_report(ytest, ypred_stk_3))
```

	precision	recall	f1-score	support
0	0.98	0.69	0.81	3517
1	0.19	0.84	0.31	304
accuracy			0.70	3821
macro avg	0.58	0.76	0.56	3821
weighted avg	0.92	0.70	0.77	3821

```
ypred_proba_stk_3=model_stk_3.predict_proba(xtest)[:,-1]
```

```
ypred_bin_stk_3=[1 if i>=0.1 else 0 for i in ypred_proba_stk_3]
```

```
print(classification_report(ytest, ypred_bin_stk_3)) #this still gives
the same result as using just 2 algorithms
```

	precision	recall	f1-score	support
0	0.99	0.46	0.63	3517
1	0.13	0.95	0.23	304

accuracy			0.50	3821
macro avg	0.56	0.71	0.43	3821
weighted avg	0.92	0.50	0.60	3821

conclusions

#Those who are leaving have a lower income, lower quarterly ratings, and are adding lower business value than those who are not

The cities where more drivers are reporting are different from the cities where the mean income is highest.

There is an opportunity for OLA to focus only on these 3 cities because business value from each driver, income for each driver, as well as driver availability are all high

Only 3 features are significant in determining driver attrition : Age, Grade, and Quarterly Rating

Keeping only these 3 features produces better results as keeping all the features

#1. Our data is imbalanced, and the use of SMOTE works well

#2. Bagging techniques like RandomForest as well as models with strong regularization features such as xgboost are not working well in our case.

#3. Boosting through Gradient Boosting or AdaBoost or LightGBM works well in predicting the recall of the target class-drivers who are leaving

#4. Increasing the complexity of the model or the number of estimators is not working out here

#our model works well when there are very few layers in the tree, when we use very few estimators, and when we lower the threshold in the binary classification model

#5. In our case, it is more important to predict whether a driver will leave, because it is more expensive to hire new drivers than retain them.

#The side effect of that is low precision, whenever in doubt, we end up predicting that a person will leave, so only 10-15% of our predictions are true that way

#Which means that our cost of incentives to existing drivers will go up in general as the cost of hiring new drivers will come down.

#However the best result is coming when we take a logistic regression model and use smote and then lower the threshold to 0.1 or lower